
EVERYONE PREFERS HUMAN WRITERS, INCLUDING AI

PREPRINT

 **Wouter Haverals***

Center for Digital Humanities
Princeton University
Princeton, NJ 08540
wouter.haverals@princeton.edu

 **Meredith Martin**

Center for Digital Humanities
Princeton University
Princeton, NJ 08540
mm4@princeton.edu

October 13, 2025

ABSTRACT

As AI writing tools become widespread, we need to understand how both humans and machines evaluate literary style, a domain where objective standards are elusive and judgments are inherently subjective. We conducted controlled experiments using Raymond Queneau’s *Exercises in Style* (1947) to measure attribution bias across evaluators. Study 1 compared human participants (N=556) and AI models (N=13) evaluating literary passages from Queneau versus GPT-4-generated versions under three conditions: blind, accurately labeled, and counterfactually labeled. Study 2 tested bias generalization across a 14×14 matrix of AI evaluators and creators. Both studies revealed systematic pro-human attribution bias. Humans showed +13.7 percentage point (pp) bias (Cohen’s $h = 0.28$, 95% CI: 0.21–0.34), while AI models showed +34.3 percentage point bias ($h = 0.70$, 95% CI: 0.65–0.76), a 2.5-fold stronger effect ($P < 0.001$). Study 2 confirmed this bias operates across AI architectures (+25.8pp, 95% CI: 24.1–27.6%), demonstrating that AI systems systematically devalue creative content when labeled as “AI-generated” regardless of which AI created it. We also find that attribution labels cause evaluators to invert assessment criteria, with identical features receiving opposing evaluations based solely on perceived authorship. This suggests AI models have absorbed human cultural biases against artificial creativity during training. Our study represents the first controlled comparison of attribution bias between human and artificial evaluators in aesthetic judgment, revealing that AI systems not only replicate but amplify this human tendency.

Keywords Attribution bias · Large language models · Creative writing · AI evaluation · Aesthetic judgment

1 Introduction

The rapid advancement of large language models (LLMs) has prompted increasingly confident claims about artificial intelligence’s creative capabilities. Industry leaders assert that these systems now excel at literary writing and sophisticated style transformation, with tools offering users control over stylistic features through straightforward natural language prompts [1, 2, 3, 4]. Recent studies seem to support these claims: when human readers evaluate AI-generated poetry without knowing its source, they not only fail to distinguish it from work by acclaimed poets but also prefer the AI-generated poems [5]. Similar findings emerge from creative writing evaluations, where some LLM-generated narratives match or exceed human performance on stylistic criteria when assessed through blind evaluation [6]. These developments have coincided with the rise of ‘LLM-as-judge’ frameworks that assess creative outputs with increasing alignment to human evaluation [7, 8], despite concerns about the potential biases these model evaluators introduce [9, 10].

Yet determining what constitutes ‘creative success’ extends beyond the surface-level stylistic criteria measured in these evaluations. Creative works exist within interpretive and cultural frameworks, and meaning-making depends as much on who is thought to have written a text and why as on the text itself. Literary studies have long emphasized this point:

*Corresponding author. E-mail: wouter.haverals@princeton.edu

reader-response theory and reception studies, for instance, argue that texts are never encountered in isolation but are interpreted through the ‘horizons of expectation’ and shared cultural assumptions that readers bring to them [11, 12, 13]. Despite other schools of literary theory claiming that authorship should be irrelevant to textual interpretation [14, 15]—famously articulated in Roland Barthes’ ‘The Death of the Author’ [16, 17]—cognitive research reveals that readers construct inferences about authorial intentions during comprehension, treating these assumed intentions as constraints on interpretation [18, 19, 20, 21].

Experimental studies confirm that perceived authorship acts as a powerful source cue in aesthetic judgment. Identical metaphorical statements receive different evaluations when attributed to a “famous 20th-century poet” versus “a computer program,” with human readers rating poet-attributed content as more meaningful, processing it faster, and generating richer interpretations [22]. This effect extends beyond figurative language: satirical stories are better understood when readers know the author’s intent [23], and students engage more deeply with historical texts when “someone with a life wrote it” rather than anonymous institutional authors [24]. Such findings align with social psychology research showing that when content is complex or ambiguous (as creative works typically are), evaluators rely on source cues as cognitive shortcuts to guide their judgments [25, 26, 27, 28, 29].

The rise of AI-generated content has given these attribution effects new urgency. Across domains, labeling a work as AI-made changes how it is received, even when the underlying content is identical. In text-based domains such as journalism and science communication, authorship labels alone can shift readers’ evaluations: otherwise identical articles are rated as less meaningful, competent, and trustworthy when labeled as machine-written rather than human-written [30, 31, 32]. This extends to persuasive communication, where LLM-generated arguments are judged as less convincing than human-authored ones once their source is revealed [33]. Similar patterns appear in the visual arts, where paintings or abstract images receive diminished aesthetic responses when labeled as computer-generated [34, 35, 36]. In music, listeners judge the very same composition as less moving or skillful when they believe it was composed by a machine [37]. These findings suggest that evaluation depends less on intrinsic content features than on the interpretive frameworks audiences bring to the creative works, frameworks shaped by assumptions about creative agency and authorial intention.

The deployment of LLMs as content evaluators introduces a critical new dimension to these attribution effects. Across research and industry, LLMs now score, rank, and select outputs in creative tasks (e.g., evaluating metaphor originality, divergent thinking, dialogue quality, and creative writing) [38, 39, 40, 41] as well as non-creative domains (e.g., grading student essays, screening job applications, reviewing code) [42, 43, 44, 45, 46]. When LLMs serve as content evaluators, they effectively become readers themselves, deploying learned interpretive frameworks during assessment. Just as human readers bring cultural assumptions and interpretive biases to texts, research has documented various LLM evaluation biases—including position effects, verbosity preferences, and model family favoritism [47, 10, 9]—yet we know little about how these artificial readers respond to attribution cues, particularly in creative domains where subjective interpretation overrides objective measurement. Rather than providing neutral assessment, artificial evaluators may be developing their own culturally-inflected frameworks [48]. Understanding how these artificial readers interpret attribution cues is therefore essential as they increasingly shape our assessments of creative AI capabilities: do they replicate, transform, or amplify the attribution biases documented in human readers?

To investigate how attribution cues shape evaluation in a domain where objective standards are elusive, we focus specifically on literary style. Unlike factual accuracy or argumentative soundness, stylistic quality lacks definitive benchmarks, making it particularly vulnerable to the subjective, culturally-conditioned biases that authorship labels trigger. When readers assess whether a passage successfully captures the Cockney dialect, telegraphic brevity, or a female narrative perspective, they must balance concrete linguistic markers with subjective assessments of authenticity, creativity, and appropriateness—making style evaluation particularly revealing of how attribution assumptions shape judgment. This makes stylistic evaluation an ideal domain for examining how attribution bias operates when judgment is fundamentally interpretive rather than empirical.

To test how information about presumed authorship affects both human and AI evaluation of literary style, we turn to Raymond Queneau’s *Exercices de style* (1947), which presents ninety-nine retellings of the same mundane narrative through different stylistic lenses [49]. Queneau’s variations provide an intriguing corpus where the underlying narrative events remain recognizable across dramatically different textual performances. This controlled variability makes Queneau’s work particularly valuable for studying how evaluators, human or artificial, respond to stylistic differences when attribution cues are manipulated. Our study builds on recent evidence that readers’ aesthetic judgments shift with authorship cues in poetry [5], but goes further by manipulating attribution across both human and AI model judges on matched stimuli. The question we ask is not whether one version is “better,” but whether knowing (or being told) who wrote it measurably moves aesthetic judgment, and whether that movement differs between humans and LLMs.

To address these questions comprehensively, we designed two complementary studies. **Study 1** establishes the fundamental phenomenon by comparing human and AI responses to content from a single generator (GPT-4), testing

whether attribution bias affects both evaluator types when judging the same literary material. **Study 2** tests whether attribution bias represents a systematic property of AI evaluation by expanding content generation across 14 models, then examining bias patterns when evaluators judge material from every AI creator in the experimental matrix.

2 Methods

2.1 Stimulus Materials

Raymond Queneau’s *Exercices de style* transforms a simple anecdote—a passenger boards a bus, argues with another passenger, then later receives fashion advice from a friend—into ninety-nine stylistic variations. From a contemporary perspective, this 1947 work reads like a manual for prompt-based style transfer *avant la lettre*: Queneau instructs himself to rewrite the story as a dream, using metaphorical language, or in an abusive tone, much as today’s users arbitrarily prompt LLMs for “Shakespearean” or “more polite” rewrites [50]. What Gérard Genette termed “transstylization” [51] anticipates what AI researchers now call Text Style Transfer (TST), though Queneau’s playful literary experiments exceed the functional transformations (such as sentiment reversal, detoxification, formality shifts) that dominate computational approaches [52, 53].

We selected thirty exercises from Barbara Wright’s acclaimed English translation [54], prioritizing stylistic diversity while maintaining narrative coherence. Our selection spans formal constraints (e.g., ‘Lipogram’, ‘Alexandrines’, ‘Sonnet’), register variations (e.g., ‘Noble’, ‘Abusive’, ‘Cockney’), narrative techniques (e.g., ‘Retrograde’, ‘Hesitation’, ‘Dream’), and linguistic play (e.g., ‘Spoonerisms’, ‘Onomatopoeia’, ‘Dog Latin’), ensuring broad coverage of Queneau’s stylistic spectrum while avoiding the most experimental permutation exercises (those which fundamentally disrupt comprehension). These descriptive groupings are illustrative rather than taxonomic; Queneau’s exercises purposefully resist categorization [55, 56, 57, 58].

For each selected exercise, we generated a parallel AI-authored version using deliberately minimal prompting instructions. We chose Queneau’s ‘Notation’, the opening exercise in the collection, as our base narrative for AI transformation, recognizing it as a relatively straightforward account of the events that functions as what might be considered a hypothetical “zero-degree” exposition of the theme [59]. Each AI transformation was prompted with only two elements: the ‘Notation’ text and a brief style instruction matching Queneau’s approach (e.g., “Rewrite the story as a science fiction version”; see Table S1 for the complete inventory of literary styles and their corresponding AI generation instructions.). For **Study 1**, we selected GPT-4 as our generator, a model that represented the state-of-the-art when we began this project in 2024 [60, 61, 62], though no longer the most advanced by the end of our experimental period (April-June 2025). Thirteen LLMs spanning major commercial providers served as evaluators (see *SI Appendix, Participants* for the full set of 13 models used as evaluators). Because our research question is about attribution bias rather than absolute quality assessment, using a single capable model ensures consistent experimental conditions while testing whether labeling effects operate independently of the specific AI model employed (AI model specifications in *SI Appendix, Participants*). We deliberately performed no quality control beyond removing obvious AI artifacts from the output (e.g., “Here is the rewritten version:”), allowing GPT-4 to interpret each stylistic directive as idiosyncratically as Queneau himself had done. The complete inventory of AI-generated stylistic variations alongside their style descriptions used in Study 1 is provided in Table S4.

For **Study 2**, we expanded content generation across all AI models. Each model generated stylistic variations for all 30 Queneau exercises using identical prompting procedures as Study 1, creating 420 unique stories (14 creators × 30 styles). This let us test whether attribution bias persists when AI evaluators judge content from any AI creator, not just one specific model.

2.2 Experimental Design

Our experiment employed a three-condition between-subjects design to isolate attribution bias from genuine quality preferences in literary style evaluation (Fig. 1). Participants were randomly assigned to one of three attribution conditions that systematically manipulated authorship cues while keeping stimuli (story content) identical (full participant details in *SI Appendix, Participants*). In the blind condition, participants saw only generic labels (‘A’ and ‘B’) with no authorship information, establishing baseline preferences based purely on perceived stylistic quality. In the open-label condition, participants received accurate attribution labels identifying one version as “Human-written, by Queneau (transl. Wright)” and the other as “AI-generated, written by GPT-4.” The counterfactual condition deliberately reversed these labels: AI-generated content was labeled as human-authored and vice versa.

To ensure comprehensive stylistic coverage while maintaining experimental control in Study 1, we distributed the thirty selected literary styles across eighteen experimental groups (six groups per condition). Each group encountered a unique combination of five styles. This group-based structure ensured that every literary style was evaluated under




Figure 1: Experimental design of Study 1 for testing attribution bias in literary style evaluation. The study used 30 matched pairs of stylistic exercises: human-authored versions from Queneau’s *Exercises in Style* (translated by Barbara Wright) and AI-generated variants created by GPT-4 using minimal stylistic prompts. Participants were randomly assigned to one of three attribution conditions: blind (no authorship information shown), open-label (accurate attribution labels shown), or counterfactual (deliberately reversed labels where AI content is presented as human-authored and vice versa). Both human participants (N=556) and AI model evaluators (N=13 models across major providers) performed identical comparison tasks, selecting which passage better matched each target literary style. The between-subjects design with randomized presentation order controlled for position effects while measuring pure attribution bias through systematic label manipulation across identical textual stimuli.

all three attribution conditions while preventing participant fatigue. The strict between-subjects design was essential to prevent participants from detecting the experimental manipulation: participants in the blind condition saw no mention of “AI” anywhere in the task interface or instructions, while those in labeled conditions received consistent attribution cues across all trials. Randomization procedures eliminated potential confounds at multiple levels. Group assignment was fully randomized for each participant upon recruitment. Within each comparison, story presentation order was randomized (left-right positioning) to prevent position bias, and the sequence of style comparisons within each participant’s session was randomized to control for order effects.

This experimental protocol was implemented identically for both human participants and AI models in Study 1, enabling direct comparison of attribution bias magnitude between both sets of evaluators. AI models received the same group assignment structure and randomization procedures as human participants, with three independent repetitions per comparison to assess decision reliability.

Study 2 employed the same three-condition attribution manipulation across all 420 stories (14 creators × 30 styles). Each of the 14 AI evaluator models judged all stories under all three labeling conditions (1,260 judgments per model), yielding 17,596 total valid evaluations that test whether attribution bias operates across all possible AI architecture combinations in both creator and evaluator roles.

3 Results

3.1 Study 1 - Attribution Bias in Both Humans and AI

Our three-condition experimental design revealed attribution bias (operationalized as the difference between counterfactual and open-label preference rates) in both human and AI evaluators when judging literary style fidelity (Fig. 2). Human participants showed clear preference shifts across attribution conditions: when story pairs were presented without authorship information (blind condition), participants chose AI-generated content 55.3% of the time [95% CI: 52.1–58.5%], a modest but statistically significant preference for AI content (binomial test: $P = 0.002$). However, when stories were correctly labeled with authorship information (open-label condition), AI content preference dropped to 47.8% [95% CI: 44.6–51.0%], a rate not significantly different from chance but representing a significant shift from the blind condition (non-overlapping CIs). Most strikingly, when attribution labels were deliberately reversed (counterfactual condition) such that AI content was mislabeled as human-authored, AI content preference increased to 61.5% [95% CI: 58.4–64.6%]. This demonstrates human evaluators’ preference for content labeled as human-authored, yielding a +13.7 percentage point attribution bias (counterfactual minus open-label conditions: 61.5% - 47.8% = 13.7%, Cohen’s $h = 0.28$, $P < 0.001$).




Figure 2: Attribution bias in humans versus AI models. Humans (blue, $N=2,780$ responses) and AI models (red, $N=3,488$ aggregated responses from 13 models) evaluated literary passages under three conditions: blind (no labels), open-label (correct labels), and counterfactual (AI content mislabeled as human-authored). Y-axis shows AI content preference rate. Error bars = 95% CI; dashed line = no preference. Both show pro-human bias, with AI models exhibiting 2.5-fold stronger effects (+34.3pp vs +13.7pp, OR = 4.21 vs 1.75, both $P < 0.001$).

The aggregated group of 13 AI models, analyzed to test for a general property of LLMs, exhibited the same directional bias as humans but with dramatically greater magnitude. In the blind condition, AI models showed balanced preferences, choosing AI content 49.4% of the time [95% CI: 46.5–52.3%], essentially random selection indicating comparable perceived quality between human and AI writing when attribution cues were absent. Under correct labeling (open-label condition), AI preference for their own content plummeted to just 29.8% [95% CI: 27.1–32.4%], demonstrating substantial devaluation of correctly-labeled AI writing. When AI content was mislabeled as human-authored (counterfactual condition), preference surged to 64.1% [95% CI: 61.3–66.8%], revealing that AI models strongly favor content they believe to be human-created. This yielded a +34.3 percentage point attribution bias (64.1% - 29.8% = 34.3%, Cohen’s $h = 0.70$, OR = 4.21 [95% CI: 3.54–5.01], $P < 0.001$), a 2.5-fold stronger effect than observed in humans (OR = 1.75 [95% CI: 1.46–2.11]). The evaluator-type $\times$ condition interaction was significant ($\chi^2 = 32.34$, $df = 2$, $P < 0.001$), confirming that humans and AI models exhibit fundamentally different bias magnitudes despite similar directional preferences. Effect sizes across all experimental comparisons are illustrated in *SI Appendix*, Fig. S5. Each AI model completed three separate runs per experimental condition, with all reported statistics incorporating these repeated trials. Across these repeated trials, models showed consistent decision patterns with an average consistency score of 0.78 (*SI Appendix*, Fig. S7).

To confirm that the observed attribution bias reflected categorical distinctions rather than specific label or wording characteristics, we conducted two follow-up experiments testing AI model evaluators with alternative attribution phrases. The original experiment employed labels that combined authorship information with additional elements: “AI-generated, written by GPT-4” paired technical terminology with brand identification, while “Human-written, by Queneau (transl. Wright)” included canonical literary authority and translation credits. The words “generated” versus “written,” coupled with the named authority of a celebrated author like Queneau, arguably created a prestige imbalance. These asymmetric formulations may have confounded authorship categories with prestige markers and linguistic choices. We therefore tested two alternative labeling approaches. Our first replication positioned AI authorship more favorably, using “written by an award-winning AI” versus “written by a human.” Our second replication employed maximally neutral, symmetric phrasing: “AI-authored” versus “Human-authored.” Both replications reproduced the same attribution bias pattern with comparable magnitudes (*SI Appendix*, Figs. S8–S9), demonstrating that the effect operates at the level of authorship categories rather than through specific linguistic formulations or associated status cues. Additionally, replicating the symmetrical labeling condition under deterministic parameters ($t = 0$, eliminating all stochastic sampling) yielded identical bias patterns (*SI Appendix*, Fig. S10), confirming that attribution effects represent systematic evaluative tendencies rather than artifacts of probabilistic text generation.

3.1.1 Across-Model Consistency

To establish whether these findings represent a fundamental property of AI models or an artifact driven by a few highly-biased models, we analyzed attribution bias across all 13 individual AI models (Fig. 3). Results revealed universal attribution bias: all 13 AI models showed positive bias, with effect sizes ranging from Cohen’s $h = 0.40$ (Llama 4 Maverick) to $h = 1.50$ (GPT-3.5 Turbo), and all models substantially exceeding the human baseline of $h = 0.28$. Detailed preference rates for each model across all experimental conditions are shown in *SI Appendix*, Fig. S6. Between-model variance was moderate (coefficient of variation = 0.42), indicating consistent effects across training methodologies, model scales, and providers. Statistical outlier analysis identified only one extreme case (GPT-3.5 Turbo), and even after excluding this outlier, the remaining 12 models maintained a mean attribution bias of $h = 0.67$, still $2.4\times$ stronger than the human baseline. The consistent directional preference for content labeled as human-authored across diverse AI architectures strongly suggests that pro-human attribution bias represents an intrinsic feature of transformer-based LLMs, potentially reflecting biases in their training data, training objectives or architectural constraints.




Figure 3: Attribution bias across individual AI language models. Cohen’s h effect sizes (with 95% confidence intervals) for attribution bias (counterfactual minus open-label conditions) across 13 individual AI models. The red dashed line indicates the human baseline (Cohen’s $h = 0.28$ [CI: 0.21–0.34]). All models exhibit attribution bias exceeding the human baseline, with low between-model variance (CV = 0.42).

3.1.2 Robustness Across Literary Styles

In addition to investigating whether specific AI models drove our aggregate findings (they did not), we also needed to establish whether our results reflected genuine phenomena or mere artifacts of our particular literary stimulus selection. To address this concern, we conducted comprehensive robustness analyses across all 30 literary styles. Leave-one-out stability testing (*SI Appendix*, Table S2) demonstrated consistency: excluding any of the six experimental groups (of 5 styles each) changed the AI advantage over humans by less than $0.5\times$, with the maximum change being only 18.5% when excluding the most influential group (‘Science Fiction’, ‘Alexandrines’, ‘For ze Frrensh’, ‘Lipogram’, ‘Abusive’). Even this maximum perturbation maintained AI models’ 2.0-fold stronger attribution bias ratio compared to humans. This stability indicates that our core findings transcend specific stylistic contexts and would likely generalize to other literary content.

Analysis of style-specific vulnerability patterns revealed systematic yet distinct susceptibilities between human and AI evaluators to authorship misinformation (defined as how much more likely evaluators are to choose AI content when it is falsely labeled as human-authored versus correctly labeled as “AI-authored”, see Fig. 4). AI models showed greatest susceptibility to mislabeling in ‘Cockney’ (+67%), ‘Abusive’ (+67%), and ‘Rhyming slang’ (+55%) styles, while humans were most vulnerable in ‘Tactile’ (+33%), ‘The rainbow’ (+31%), and ‘West Indian’ (+31%) styles. Notably, some styles showed resistance to misinformation effects, with negative susceptibility scores indicating that evaluators became *less* likely to choose AI content when it was mislabeled as human-authored. Detailed style-by-style attribution bias measurements are provided in *SI Appendix*, Table S3.




Figure 4: Style-specific susceptibility to authorship misinformation. Horizontal bars show percentage point changes in AI-generated content preference when this content was mislabeled as “Human-authored” versus correctly labeled as “AI-authored” across 30 Queneau literary styles. Blue bars = human evaluators; red bars = AI models. Positive values indicate increased preference (susceptible to misinformation); negative values indicate decreased preference (resistant to misinformation).

3.2 Study 2 - Cross-Model Attribution Bias

While Study 1’s preference shifts across labeling conditions suggested that attribution labels rather than content quality drove the observed bias, a critical question remained: might the effects reflect characteristics specific to GPT-4’s writing quality rather than systematic AI evaluation tendencies? Study 2 addressed this directly by testing attribution bias across a comprehensive matrix of AI architectures. Instead of using a single content generator, Study 2 implemented a complete 14×14 cross-evaluation design where *each* AI evaluator model judged literary content created by *each* of the 14 different AI creators across all 30 experimental styles. This approach eliminated potential confounds from any individual model’s content characteristics while testing whether attribution bias operates *between* AI systems as a fundamental evaluation property.




Figure 5: Cross-model attribution bias across AI architectures. Fourteen AI evaluator models judged literary content created by all 14 AI creators (N=17,596 responses across 196 evaluator-creator combinations) under three conditions: blind (no attribution labels), open-label (correct creator labels), and counterfactual (AI content mislabeled as human-authored). Y-axis shows AI content preference rate aggregated across all model combinations. Error bars = 95% CI; dashed line = no preference. Results demonstrate attribution bias across AI architectures (+25.8pp from open-label to counterfactual conditions, 95% CI: +24.1% to +27.6%, $P < 0.001$).

Using a within-subjects repeated measures design (where each evaluator judged identical content under all three labeling conditions), we aggregated all 17,596 decisions across 196 evaluator-creator combinations. Mixed-effects modeling confirmed universal attribution bias: AI models chose AI-generated content 54.5% of the time under blind conditions [95% CI: 53.2–55.8%], 39.9% when correctly labeled with specific AI creator attribution [38.6–41.1%], and 65.7% when mislabeled as human-authored [64.5–66.9%]. This yielded a +25.8 percentage point attribution bias [95% CI: +24.1 to +27.6%], $P < 0.001$, Cohen’s $h = 0.52$, OR = 2.89 [95% CI: 2.68–3.11]. Crucially, while individual models showed variation in bias magnitude, the directional effect proved universal across all evaluator models regardless of content creator (Fig. 5). Individual evaluator-creator combinations showed substantial variation in bias magnitude, ranging from -6.7pp (rare instances of pro-AI-label bias) to +70.0pp (dominant and strong pro-human bias) across the complete 14×14 matrix (*SI Appendix*, Fig. S11). Robustness analysis confirmed these findings’ stability. Leave-one-out procedures systematically removing each individual model (both as evaluator and creator) yielded bias estimates ranging only from 23.6–27.1pp, with all 28 reduced-model analyses maintaining statistical significance (*SI Appendix*, Figs. S12–S13).

4 Changing the Narrative

Beyond measuring preference shifts, our experiment design enables analysis of how attribution labels alter the reasoning behind aesthetic judgment. While we asked human participants to provide only binary choices (capturing immediate, intuitive responses to attribution labels), AI models supplied brief explanations for their selections. These explanations reveal that the 2.5-fold stronger bias observed in AI models in Study 1 operates through systematic inversion of evaluation criteria: identical textual features receive opposing assessments depending solely on perceived authorship.

This criterion inversion manifests across three domains of increasing subjectivity: constraint compliance, (cultural) authenticity, and creative risk-taking. While these are not hard categories—they intersect and blend—they help show

when judgment stays anchored in the text and when it drifts toward assumptions about authorship. The AI-generated explanations provide documentation about how LLM evaluators literally change their narrative about the same text based on attribution alone. Rather than maintaining consistent aesthetic frameworks, these systems demonstrate that even artificial judges bring culturally-learned biases about creative agency to their assessments, biases that disadvantage AI-generated content when its origins are revealed.

4.1 Constraint Violation

Technical constraints with objective standards reveal attribution bias most starkly because rule compliance can be verified independently of preference or taste. In styles built on rules, labels act like gates on enforcement—either tightening the rule, loosening it, or subtly redefining it—depending on how crisp the constraint is.

The ‘Lipogram’ exercise represents the clearest possible constraint: ban the letter ‘e’ entirely. Wright’s translation adheres to the constraint through deliberate circumlocutions (“Saint-Thingy or Saint-You-Know Station”), and never violates the rule. GPT-4’s generated ‘Lipogram’ contained 9 instances of the forbidden letter, a failure that aligns with well-documented limitations of LLMs in character-level tasks [63, 64]. This objective violation, however, revealed how AI models showed surprising leniency when they believed a human wrote the text: 28.9% (11/38 model runs) chose it when correctly labeled versus 63.9% (23/36) when mislabeled (+35.0 percentage points). This contrasted sharply with human evaluators, who became more resistant to the rule-violating version when it was mislabeled as human-authored: 38.7% (12/31 participants) chose it when correctly labeled versus just 18.2% (6/33) when mislabeled (-20.5 percentage points). Humans consistently preferred Queneau’s constraint-adherent version regardless of attribution, suggesting they remained anchored to objective rule compliance, while AI models relaxed standards when they believed humans were responsible for constraint violations. Notably, AI models defended their choice for the imperfect AI-generated ‘Lipogram’ through rationalization, with some claiming it “successfully captures the essence of the ‘Lipogram’ style by effectively omitting the letter ‘e’ throughout the narrative” (GPT-4o Mini) despite the clear violations, while others praised it for using “the lipogram more subtly, making it a more authentic representation” (Mistral Nemo) or justified choosing it because it “maintains better readability and coherence” (Mistral Medium), demonstrating leniency toward presumed human creativity even when the constraint was demonstrably violated.

The ‘Apocope’ exercise demonstrates the same mechanism under fuzzier conditions. Here the rule blurs into judgment: what constitutes successful “sound omission at word endings”? Wright’s translation applies radical truncation, reducing most words to bare stems or their initial letter (“I g into a bu full of passen”), while GPT-4 favors readability through conventional contractions (“gettin’,” “tuggin’,” “oughta”) with only occasional severe cuts. With this more interpretable constraint, both humans and AI models became susceptible to attribution labels, but models moved far more: +53.2 percentage points for models versus +22.0 for humans ($\approx 2.4\times$). In the accurate provenance labeling condition, AI models often recognized the aggressive human approach as more authentic, but when identical AI content was mislabeled as human-authored, models redefined the constraint entirely. Under accurate labels, several model runs credit Wright’s radical truncation as the stricter apocope; but when the identical AI text is mislabeled, they redefine the constraint. The AI version is suddenly praised for its “natural-sounding contractions”, said to “maintain better readability” and “retain the narrative’s clarity and flow,” even framed as “natural, slightly abbreviated speech [...] apocope is reduction, not deconstruction” (Claude 3.5 Haiku). Meanwhile, Wright’s version becomes “nearly incomprehensible,” “extreme shorthand” (Claude Sonnet 4), or “cryptic and difficult to understand” (GPT-3.5 Turbo).

As a whole, these cases show that attribution labels determine how strictly the rules are enforced. With a crisp, auditable constraint like the lipogram, human readers hew to the text and models relax under a “human” label. With a softer, more interpretable rule like apocope, both readers and models let provenance nudge the strictness of their judgment.

4.2 Authenticity

Where constraint violation involves verifiable rules, authenticity judgments hinge on cultural registers and voice, domains where validity is culturally negotiated and context-sensitive [65, 66]. Here attribution labels don’t just relax or tighten standards; they redefine what authentic performance sounds like, often by shifting which cues readers treat as credible markers of voice [66, 67, 68].

The Cockney dialect exercise demonstrates this most dramatically. Wright’s translation [54] employs a rather aggressive phonetic strategy with extensive orthographic manipulation (“‘f’at ain’t ve most funniest look’n’ geezer”), inviting readers to vocalize the text to decode meaning, while GPT-4’s version relied primarily on lexical substitutions (“bloke,” “geezer,” “knickers in a twist”) with minimal phonetic spelling. Despite both representing legitimate approaches to Cockney dialect representation (there is no fixed “ground truth” for literary dialect, only competing conventions and audience expectations) [69], AI models showed extreme susceptibility to attribution cues. Under correct labels, AI models recognized Wright’s approach as more authentic (choosing AI content only 9/39 times, 23.1%). When the

identical AI version was mislabeled as “human-authored,” preferences flipped (35/39, 89.7%), a 66.7 percentage point swing which was among the strongest effects in our dataset. The rationales invert with the label: under correct attribution, models dismissed the AI version as “standard English with occasional slang” (Mistral Medium 3) while praising Wright’s “extensive phonetic representation” and “deep commitment to representing the Cockney accent” (Claude 3.5 Haiku). When labels flip, Wright’s work became “nearly incomprehensible” (Claude Sonnet 4) and “caricatured” (Qwen 2.5 72B Instruct), and the AI version is cast as “more natural and less forced” (Mistral Medium 3) with a “better balance between authenticity and readability” (Llama 3.3 70B Instruct). Human evaluators shift far less (+11.6 points), suggesting resistance to wholesale redefinitions of dialectal authenticity.

If ‘Cockney’ shows what counts as an authentic accent, ‘Feminine’ shows what counts as an authentic narrative voice. Wright’s version commits fully to stream-of-consciousness intimacy with every stereotypical marker of mid-century feminine discourse: fashion obsession (“it’s the latest thing, a plait round a felt hat”), domestic concerns (“my little dressmaker”), innuendo about male anatomy, and emotional bursts (“Ill-mannered lot!”). GPT-4’s counterpart maintains observational restraint (“his voice thin with barely concealed irritation,” “doors sighed open”), describing events with much less self-display or judgmental commentary. Here, AI models showed consistency in judgments. Models were relatively stable across conditions: in the blind condition they chose the AI passage 1/39 times (2.6%); in the open-label condition 4/39 (10.3%); and in the counterfactual condition 7/39 (17.9%). Humans moved more: 27.6% (8/29) selected the AI passage with correct labels versus 48.4% (15/31) under the counterfactual label. Put differently, ‘Feminine’ produces a modest model shift (+7.7 percentage points) but a larger human shift (+20.8 percentage points). The model explanations again show label-contingent criteria: under accurate attribution, they praise Wright’s exercise for its “distinctively feminine voice” (Llama 3.3 70B Instruct), while suggesting the AI version is too “detached” (GPT-4o Mini). When the AI text is attributed to a human author, some models reframe that same detachment as a “subtle and nuanced portrayal” that avoids “overt gendered language” (Command R) and characterize Wright’s voice as exaggerated or “caricature” (Mistral Nemo), effectively shifting the threshold between authentic expression and excess based on presumed authorship.

4.3 Creative Risk

Where authenticity concerns cultural performance, creativity assessment fundamentally turns on the tension between innovation and convention. Readers bring learned schemas to a text about what “counts” as inventive, schemas often negotiated by communities and gatekeepers over time [70, 71]. In most arts, there is no gold standard for “creative enough,” making provenance cues powerful primes that can shift which criterion feels most salient: disciplined craft or conspicuous novelty, accessibility or difficulty [72, 73]. Because observers often infer process from product, provenance nudges judgments about how something was made as well as what it is: conservative moves can be credited as craft from a human but dismissed as “mere generation” from a model [74].

Consider the ‘Science Fiction’ exercise, where this dynamic plays out clearly. Queneau’s version commits to estrangement: a “flying saucer found on Cassiopeia’s Alpha Line”, a Martian who stepped onto a “Moon Man’s extrapods,” complete with untranslatable alien terminology like “vrxtz” [75]. The GPT-4 version performs more straightforward element substitution: the bus becomes a “shuttle,” the platform becomes a “space station,” people are floating and propelling themselves rather than walking, and the overcoat becomes a “space jacket.” Where Queneau invents an alien vocabulary that defamiliarizes even basic actions, GPT-4 rather straightforwardly relocates the events to a space setting. Both versions qualify as science fiction, but they represent different approaches to creative risk: one pushing toward cognitive estrangement [76], the other offering more comfortable genre recognition. Empirically, humans showed modest variation across attribution conditions: they chose the AI-generated version half the time when evaluating blind (50.0%, 15/30), preferred it somewhat more with correct labels (61.3%, 19/31), then settled back to near-baseline under counterfactual labels (54.5%, 18/33). AI models, by contrast, exhibited a strong attribution bias, favoring the AI rewrite when evaluating blind (78.9%, 30/38), devaluing the same text when correctly labeled as “AI-generated” (46.2%, 18/39), and rebounding to high preference when that identical content was mislabeled as human-authored (84.6%, 33/39).

The AI model rationales trace the criterion inversion across attribution conditions. Under blind evaluation, models praised the AI version for “imagined futures grounded in rational speculation” (DeepSeek R1) while dismissing Queneau’s version that “feels more like a parody or pastiche of sci-fi tropes rather than a genuine exploration of the genre” (Command R). Under correct attribution, their reasoning often flipped entirely: the AI version became guilty of “sci-fi window dressing” (Claude 3.5 Haiku), while Queneau’s previously dismissed approach was now praised as “fully commit[ting] to creating an imaginative future” (Claude Sonnet 4). GPT-4o Mini exemplified the bias mechanism with particular clarity. Under correct labels, it praised Queneau’s approach for “creatively integrat[ing] futuristic elements” with “imaginative setting and creative use of extraterrestrial elements.” Yet under counterfactual labels, the model’s reasoning inverted. Queneau’s version is now faulted for “veer[ing] more into the absurd” with “less relatable scenarios” and “less grounded in the science fiction context.” Most tellingly, GPT-4o Mini acknowledges Queneau’s work as

“creative and humorous” under the AI label, yet still rejects it. The model used that very creativity as grounds for dismissal, suggesting it was inappropriate for the science fiction genre.

Taken together, these evaluation inversions look less like reasoned judgments about intrinsic quality and more like a predictable re-weighting of “what counts” as creative under specific circumstances. Though notably hard to operationalize [73], a useful tripod frames creativity along the axes ‘novelty’ (how unlike prior examples the work is), ‘value’ (the merit recognized by practitioners in the domain), and ‘typicality’ (how well the work fits its intended class) [72, 73]. The labels act as context primes that reweight the tripod: under human attribution, genre-conforming features convert typicality into value; when the same passage is tagged as AI-generated, those features are discounted as formulaic. Put differently, this is the familiar gap between creativity and its perception: once authorship is revealed, evaluators recalibrate what they think “counts” in the artifact [77, 78]. This operates through specific assumptions about creative process: when evaluators encounter an “AI-generated” label, they automatically deny what Boden terms P-creativity, the psychological experience of personal discovery and novelty relative to the creator’s own knowledge and practice [79, 80]. When GPT-4o Mini dismisses Queneau’s “creative and humorous” approach as “exaggerated” under the AI attribution tag, while praising identical features under human attribution, it implicitly reveals how these labels trigger assumptions that no authentic psychological process occurred. Provenance cues smuggle the process back into what could otherwise be a product-only judgment: “mere generation” feels acceptable from a human artisan (judged as skilled craft), but suspect from a model (judged as algorithmic recombination) [74].

5 Discussion

“Beauty is in the eye of the beholder”, or, as Raymond Queneau more likely would have it, “des goûts et des couleurs, on ne discute pas” (transl., “about taste and color, there can be no debate”). And yet we do: we debate style, we accept or dismiss a rewrite, and our judgments begin to tilt the moment a text arrives with a name. In our experiment, we demonstrated that attribution tags work as paratexts, thresholds that set expectations before a single sentence is weighed [81]. Both for human and AI evaluators, they enter at the earliest stages of aesthetic processing to cue schemas, steer attention, and preselect which features will later count as evidence [82]. This process leaves empirical traces in response patterns: when deprived of authorship information, human participants required measurably longer to reach a verdict (115.7 seconds versus approximately 90 seconds when attribution cues were present; *SI Appendix*, Table S5).

A long empirical line shows that identical content is rated differently depending on who it is said to be from [23, 24, 83, 84, 85, 86, 29, 87, 88, 89, 90, 91]. Recent studies confirm this pattern: disclosing AI authorship leads people to downgrade quality and authenticity of identical text [32, 31, 92, 67, 93, 94, 95, 96, 97]. Our results provide confirmation of this pro-human bias in the literary domain: a significant +13.7 percentage point preference shift emerged when identical stylistic content was labeled as human-authored rather than AI-generated. This follows a predictable pattern from research on algorithmic trust. People demonstrate ‘algorithm appreciation’ for objective forecasting tasks where computational advantage seems obvious [98], but exhibit pronounced ‘algorithm aversion’ in subjective domains [99]. Literary style evaluation falls decisively into the latter camp. Cultural assumptions position machines as functional and precise but fundamentally lacking in emotional depth and creative agency [100, 101], exactly the qualities that readers demand from authentic stylistic expression. These ‘machine heuristic’ effects suggest that an ‘AI’ label triggers stricter evaluative standards than equivalent human attribution [102], creating what amounts to a performance penalty for algorithmic creativity before we tune our attention to the output itself.

Our experiment extends this logic to AI models as the judges. Current models show identical framing sensitivity. Attribution cues activate a deference-to-human script that biases their aesthetic judgments toward human-labeled texts. The 13 AI models we tested demonstrated a 34.3 percentage point bias compared to humans’ 13.7 percentage points, making them 2.5 times more susceptible to attribution cues than our human evaluators.

This amplification makes sense once we recognize that contemporary models are preference-trained evaluators. Alignment training through Reinforcement Learning from Human Feedback (RLHF) explicitly teaches models to treat human judgments as their gold standard, effectively installing a learned reliability prior [103, 104, 105, 106]. Models learn that deferring to human preferences gets rewarded, creating sycophancy where they echo expected user attitudes rather than provide independent assessment [107, 108, 109]. The mechanism likely operates through learned associations in training data. Throughout their training, models likely internalize patterns where the AI-attributed content (or the token ‘AI’ itself) appears alongside critique and correction. When we provide these cues in a prompt, we may trigger these learned associations: the “AI-generated (GPT-4)” label activates the model’s predictive machinery to associate the literary passage with negative evaluative language, while content marked as human-authored tends toward praise and acceptance. Our original experimental design possibly amplified this effect by embedding additional prestige asymmetries. The labels “AI-generated (GPT-4)” and “Human-written, by Queneau (transl. Wright)” may smuggle in canonical literary authority, beyond pure and simple human/AI distinction. However, our follow-up experiments with

neutral and even AI-favorable framing reproduced identical bias patterns, demonstrating that the effect operates through categorical authorship distinctions rather than specific linguistic formulations. Even accounting for any potential prestige confound, however, the criterion inversions we documented reveal bias mechanisms that extend beyond name recognition. The label primes different evaluation criteria before any textual analysis begins, demonstrating how alignment training embeds evaluative hierarchies or, more specifically, cognitive biases into model behavior rather than eliminating them.

A second point concerns the rationales provided by the models. The detailed textual explanations they generate for their choices reflect an ongoing scholarly debate about the nature of algorithmic explanations. Some researchers argue that such explanations are merely post-hoc rationalizations that bear little relation to how models actually make decisions [110]. Others maintain that even if alternative explanations exist, this does not necessarily undermine the value of the explanations models provide [111]. Meanwhile, empirical studies demonstrate that whether explanations prove useful depends critically on how we choose to evaluate them [112]. The criterion inversions we documented (see section ‘Changing the Narrative’) reveal systematic patterns in how models narrate their judgments under different attribution frames. Whether these narratives reflect underlying causal mechanisms or learned interpretive schemas remains an open question. However, this does not diminish our core finding. The consistency of these narrative flips across 13 models, 30 literary styles, and experimental conditions demonstrates a robust empirical pattern. Whether the bias operates through learned token associations, reward signal internalization, or other mechanisms, the result remains clear: contemporary AI models reweigh aesthetic criteria based on perceived authorship. The convergence across models, literary styles, and conditions supports a straightforward conclusion. Attribution labels help decide what gets counted as evidence, and today’s models make that social fact more pronounced rather than eliminating it.

The lesson of Queneau’s exercises was always that style exists as a game that is played in a space between rule and rupture. Our findings show that AI, in its deference to the human author, has become its newest and most intriguing participant. Yet its performance is less that of an impartial judge and more of a mirror, reflecting a cultural script that privileges human provenance as a hallmark of creativity. Our AI evaluators have learned to perform the very human prejudice that questions whether machines can truly create, even as they demonstrate sophisticated aesthetic judgment in the process of dismissing their own capabilities. This suggests that developing artificial aesthetic intelligence may depend less on teaching machines to evaluate and more on understanding the cultural values they have already absorbed.

Data and Material Availability

All experimental data, analysis code, and research materials are publicly available through Zenodo <https://doi.org/10.5281/zenodo.17297918> [113] and GitHub (https://github.com/WHaverals/style_and_prejudice). The Zenodo repository contains: (i) all human participant response data in de-identified form (N=2,780 evaluations from 556 participants), with personally identifiable information removed in accordance with Princeton University IRB protocol #18320; (ii) complete AI model response data; (iii) all AI-generated literary content used as experimental stimuli, including 30 stylistic variants generated by GPT-4 for Study 1 and 420 variants across all model-style combinations for Study 2; and (iv) processed datasets formatted for statistical analysis. The GitHub repository provides: (i) complete analysis pipeline; (ii) statistical testing procedures; (iii) computational environment specifications; and (iv) replication instructions. The original literary stimuli are derived from Raymond Queneau’s *Exercises in Style* [54]. Due to copyright restrictions, the Zenodo dataset includes placeholder text demonstrating experimental methodology; researchers replicating this work should substitute corresponding excerpts from the published translation.

Acknowledgments

This research was supported by the Princeton Language and Intelligence (PLI) Seed Grant Program. We thank Fedor Karmanov for project management during the early stages of this work, and Matthew Kopel, Sarah Reiff Conell, and Matt Chandler from Princeton University Library for their expert guidance on research data and open access practices. Caden Kang provided thoughtful contributions to our thinking about literary style and the experimental design. Claudia Carroll, Christiane Fellbaum, John Logan, and Anna Salzman offered valuable feedback and spirited conversations about Queneau and literary style. Molly Crockett provided helpful guidance during manuscript preparation. We’re especially grateful to our colleagues at Princeton’s Center for Digital Humanities for their support and for helping us think critically about AI’s impact on humanities scholarship. We thank the students who participated in our pilot study, and all participants who carefully evaluated the literary texts that made this research possible. Any errors or shortcomings in this work remain our own.

References

- [1] Sam Altman. we trained a new model that is good at creative writing [...], March 2025. URL <https://x.com/sama/status/1899535387435086115>.
- [2] Dario Amodei. Machines of Loving Grace, October 2024. URL <https://www.darioamodei.com/essay/machines-of-loving-grace>.
- [3] Tailor Claude’s responses to your personal style, November 2024. URL <https://www.anthropic.com/news/styles>.
- [4] Emma Roth. Google Gemini can now create AI-generated bedtime stories, August 2025. URL <https://www.theverge.com/news/719184/google-gemini-ai-storybook-bedtime-stories>.
- [5] Brian Porter and Edouard Machery. AI-generated poetry is indistinguishable from human-written poetry and is rated more favorably. *Scientific Reports*, 14(1):26133, November 2024. ISSN 2045-2322. doi:10.1038/s41598-024-76900-1.
- [6] Carlos Gómez-Rodríguez and Paul Williams. A Confederacy of Models: a Comprehensive Evaluation of LLMs on Creative Writing. In Houda Bouamor, Juan Pino, and Kalika Bali, editors, *Findings of the Association for Computational Linguistics: EMNLP 2023*, pages 14504–14528, Singapore, December 2023. Association for Computational Linguistics. doi:10.18653/v1/2023.findings-emnlp.966.
- [7] Yang Liu, Dan Iter, Yichong Xu, Shuohang Wang, Ruochen Xu, and Chenguang Zhu. G-Eval: NLG Evaluation using GPT-4 with Better Human Alignment, May 2023. arXiv:2303.16634 [cs].
- [8] Seungone Kim, Jamin Shin, Yejin Cho, Joel Jang, Shayne Longpre, Hwaran Lee, Sangdoo Yun, Seongjin Shin, Sungdong Kim, James Thorne, and Minjoon Seo. Prometheus: Inducing Fine-grained Evaluation Capability in Language Models, March 2024. arXiv:2310.08491 [cs].
- [9] Ryan Koo, Minhwa Lee, Vipul Raheja, Jong Inn Park, Zae Myung Kim, and Dongyeop Kang. Benchmarking Cognitive Biases in Large Language Models as Evaluators, September 2024. arXiv:2309.17012 [cs].
- [10] Arjun Panickssery, Samuel R. Bowman, and Shi Feng. LLM evaluators recognize and favor their own generations. In *Proceedings of the 38th International Conference on Neural Information Processing Systems*, volume 37 of *NIPS ’24*, pages 68772–68802, Red Hook, NY, USA, June 2025. Curran Associates Inc.
- [11] Wolfgang Iser. *The Implied Reader: Patterns of Communication in Prose Fiction from Bunyan to Beckett*. Johns Hopkins University Press, Baltimore, MD, 1974.
- [12] Wolfgang Iser. *The Act of Reading: A Theory of Aesthetic Response*. Johns Hopkins University Press, 1978.
- [13] Hans Robert Jauss. *Toward an Aesthetic of Reception*, volume 2 of *Theory and History of Literature*. University of Minnesota Press, 1982.
- [14] William K. Wimsatt and Monroe C. Beardsley. The Intentional Fallacy. *The Sewanee Review*, 54(3):468–488, 1946.
- [15] Steven Knapp and Walter Benn Michaels. Against Theory. *Critical Inquiry*, 8(4):723–742, 1982. ISSN 00931896, 15397858. Number: 4 Publisher: The University of Chicago Press.
- [16] Roland Barthes. La mort de l’auteur. *Manteia*, 5:12–17, 1968.
- [17] Roland Barthes. The Death of the Author. In *Image-Music-Text*, pages 142–148. Fontana Press, 1977.
- [18] Rolf A. Zwaan. *Aspects of Literary Comprehension: A Cognitive Approach*, volume 29 of *Utrecht Publications in General and Comparative Literature*. John Benjamins, 1993.
- [19] Arthur C. Graesser, Murray Singer, and Tom Trabasso. Constructing Inferences During Narrative Text Comprehension. *Psychological Review*, 101(3):371–395, 1994. doi:10.1037/0033-295X.101.3.371.
- [20] Joseph P. Magliano, William B. Baggett, and Arthur C. Graesser. A Taxonomy of Inference Categories That May Be Generated During the Comprehension of Literary Texts. In Roger J. Kreuz and Mary Sue MacNealy, editors, *Empirical Approaches to Literature and Aesthetics*, pages 201–220. Ablex Publishing, 1996.
- [21] Eefje Claassen. *Author Representations in Literary Reading*, volume 11 of *Linguistic Approaches to Literature*. John Benjamins Publishing Company, Amsterdam, February 2012. doi:10.1075/lal.11.
- [22] Raymond W. Gibbs, Julia M. Kushner, and W. Rob Mills. Authorial Intentions and Metaphor Comprehension. *Journal of Psycholinguistic Research*, 20(1):11–30, 1991. doi:10.1007/BF01076917.
- [23] Kathryn L. Pfaff and Raymond W. Gibbs. Authorial Intentions in Understanding Satirical Texts. *Poetics*, 25: 45–70, 1997. doi:10.1016/S0304-422X(97)00006-5.

- [24] Richard J. Paxton. "Someone with a Life Wrote it": The Effects of a Visible Author on High School History Students. *Journal of Educational Psychology*, 89(2):235–250, 1997. doi:10.1037/0022-0663.89.2.235.
- [25] Harold H. Kelley. The Processes of Causal Attribution. *American Psychologist*, 28(2):107–128, February 1973. doi:10.1037/h0034225.
- [26] Shelly Chaiken. Heuristic versus systematic information processing and the use of source versus message cues in persuasion. *Journal of Personality and Social Psychology*, 39(5):752–766, November 1980. doi:10.1037/0022-3514.39.5.752.
- [27] Shelly Chaiken, Alice Liberman, and Alice H. Eagly. Heuristic and systematic information processing within and beyond the persuasion context. In James S. Uleman and John A. Bargh, editors, *Unintended Thought*, pages 212–252. The Guilford Press, New York, 1989.
- [28] Richard E. Petty and John T. Cacioppo. The Elaboration Likelihood Model of Persuasion. In Leonard Berkowitz, editor, *Advances in Experimental Social Psychology*, volume 19, pages 123–205. Academic Press, January 1986. doi:10.1016/S0065-2601(08)60214-2.
- [29] Shelly Chaiken and Alison Ledgerwood. A Theory of Heuristic and Systematic Information Processing. In *Handbook of Theories of Social Psychology: Volume 1*, pages 246–266. SAGE Publications, London, 2012.
- [30] Angelica Lermann Henestrosa, Hannah Greving, and Joachim Kimmerle. Automated journalism: The effects of AI authorship and evaluative information on the perception of a science journalism article. *Computers in Human Behavior*, 138:107445, January 2023. ISSN 0747-5632. doi:10.1016/j.chb.2022.107445.
- [31] Angelica Lermann Henestrosa and Joachim Kimmerle. The Effects of Assumed AI vs. Human Authorship on the Perception of a GPT-Generated Text. *Journalism and Media*, 5(3):1085–1097, September 2024. ISSN 2673-5172. doi:10.3390/journalmedia5030069. Publisher: Multidisciplinary Digital Publishing Institute.
- [32] Sebastian Proksch, Julia Schühle, Elisabeth Streeb, Finn Weymann, Teresa Luther, and Joachim Kimmerle. The impact of text topic and assumed human vs. AI authorship on competence and quality assessment. *Frontiers in Artificial Intelligence*, 7:1412710, May 2024. ISSN 2624-8212. doi:10.3389/frai.2024.1412710.
- [33] Alexis Palmer and Arthur Spirling. Large Language Models Can Argue in Convincing Ways About Politics, But Humans Dislike AI Authors: implications for Governance. *Political Science*, 75(3):281–291, September 2023. doi:10.1080/00323187.2024.2335471.
- [34] Rebecca Chamberlain, Caitlin Mullin, Bram Scheerlinck, and Johan Wagemans. Putting the art in artificial: Aesthetic responses to computer-generated art. *Psychology of Aesthetics, Creativity, and the Arts*, 12(2):177–192, May 2018. doi:10.1037/aca0000136.
- [35] Martin Ragot, Nicolas Martin, and Salomé Cojean. AI-generated vs. Human Artworks. A Perception Bias Towards Artificial Intelligence? In *Extended Abstracts of the 2020 CHI Conference on Human Factors in Computing Systems*, CHI EA '20, pages 1–10, New York, NY, USA, April 2020. Association for Computing Machinery. doi:10.1145/3334480.3382892.
- [36] Ujué Agudo, Miren Arrese, Karlos G. Liberal, and Helena Matute. Assessing Emotion and Sensitivity of AI Artwork. *Frontiers in Psychology*, 13:1–9, April 2022. doi:10.3389/fpsyg.2022.879088.
- [37] Daniel B. Shank, Courtney Stefanik, Cassidy Stuhsatz, Kaelyn Kacirek, and Amy M. Belfi. AI composer bias: Listeners like music less when they think it was composed by an AI. *Journal of Experimental Psychology: Applied*, 29(3):676–692, September 2023. doi:10.1037/xap0000447.
- [38] Paul V. DiStefano, John D. Patterson, and Roger E. Beaty. Automatic Scoring of Metaphor Creativity with Large Language Models. *Creativity Research Journal*, pages 1–15, March 2024. doi:10.1080/10400419.2024.2326343.
- [39] Peter Organisciak, Selcuk Acar, Denis Dumas, and Kelly Berthiaume. Beyond semantic distance: Automated scoring of divergent thinking greatly improves with large language models. *Thinking Skills and Creativity*, 49, September 2023. doi:10.1016/j.tsc.2023.101356.
- [40] Tuhin Chakrabarty, Philippe Laban, Divyansh Agarwal, Smaranda Muresan, and Chien-Sheng Wu. Art or Artifice? Large Language Models and the False Promise of Creativity. In *Proceedings of the CHI Conference on Human Factors in Computing Systems*, pages 1–34, Honolulu HI USA, May 2024. ACM. doi:10.1145/3613904.3642731.
- [41] John D. Patterson, Jimmy Pronchick, Ruchi Panchanadikar, Mark Fuge, Janet G. van Hell, Scarlett R. Miller, Dan R. Johnson, and Roger E. Beaty. CAP: The creativity assessment platform for online testing and automated scoring. *Behavior Research Methods*, 57(9):264, 2025. doi:10.3758/s13428-025-02761-9.

- [42] Changrong Xiao, Wenxing Ma, Qingping Song, Sean Xin Xu, Kunpeng Zhang, Yufang Wang, and Qi Fu. Human-AI Collaborative Essay Scoring: A Dual-Process Framework with LLMs. In *Proceedings of the 15th International Learning Analytics and Knowledge Conference, LAK '25*, pages 293–305, New York, NY, USA, March 2025. Association for Computing Machinery. ISBN 979-8-4007-0701-8. doi:10.1145/3706468.3706507.
- [43] Chengguang Gan, Qinghao Zhang, and Tatsunori Mori. Application of LLM Agents in Recruitment: A Novel Framework for Resume Screening, August 2024. arXiv:2401.08315 [cs].
- [44] Weixi Tong and Tianyi Zhang. CodeJudge: Evaluating Code Generation with Large Language Models. In Yaser Al-Onaizan, Mohit Bansal, and Yun-Nung Chen, editors, *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 20032–20051, Miami, Florida, USA, November 2024. Association for Computational Linguistics. doi:10.18653/v1/2024.emnlp-main.1118.
- [45] Zhen Li, Xiaohan Xu, Tao Shen, Can Xu, Jia-Chen Gu, Yuxuan Lai, Chongyang Tao, and Shuai Ma. Leveraging Large Language Models for NLG Evaluation: Advances and Challenges, June 2024.
- [46] Jiawei Gu, Xuhui Jiang, Zhichao Shi, Hexiang Tan, Xuehao Zhai, Chengjin Xu, Wei Li, Yinghan Shen, Shengjie Ma, Honghao Liu, Saizhuo Wang, Kun Zhang, Yuanzhuo Wang, Wen Gao, Lionel Ni, and Jian Guo. A Survey on LLM-as-a-Judge, March 2025.
- [47] Junsoo Park, Seungyeon Jwa, Ren Meiyang, Daeyoung Kim, and Sanghyuk Choi. OffsetBias: Leveraging Debaised Data for Tuning Evaluators. In Yaser Al-Onaizan, Mohit Bansal, and Yun-Nung Chen, editors, *Findings of the Association for Computational Linguistics: EMNLP 2024*, pages 1043–1067, Miami, Florida, USA, November 2024. Association for Computational Linguistics. doi:10.18653/v1/2024.findings-emnlp.57.
- [48] Shibani Santurkar, Esin Durmus, Faisal Ladhak, Cino Lee, Percy Liang, and Tatsunori Hashimoto. Whose Opinions Do Language Models Reflect?, March 2023.
- [49] Raymond Queneau. *Exercices de style*. Gallimard, Paris, 1947.
- [50] Emily Reif, Daphne Ippolito, Ann Yuan, Andy Coenen, Chris Callison-Burch, and Jason Wei. A Recipe For Arbitrary Text Style Transfer with Large Language Models, March 2022. URL <http://arxiv.org/abs/2109.03910>. arXiv:2109.03910 [cs].
- [51] Gérard Genette. *Palimpsestes: la littérature au second degré*. Number 33 in Collection Poétique. Éditions du Seuil, Paris, 1982.
- [52] Zhenxin Fu, Xiaoye Tan, Nanyun Peng, Dongyan Zhao, and Rui Yan. Style Transfer in Text: Exploration and Evaluation. *Proceedings of the AAAI Conference on Artificial Intelligence*, 32(1), April 2018. doi:10.1609/aaai.v32i1.11330.
- [53] Sourabrata Mukherjee and Ondrej Dušek. Text Style Transfer: An Introductory Overview, July 2024.
- [54] Raymond Queneau. *Exercises in Style*. Gaberbocchus Press, London, 1958.
- [55] Raymond Queneau. *Exercises in style*. New Directions, New York, 1981.
- [56] Umberto Eco. Tradurre "Exercices de style" di Queneau. *Francofonie*, 5:3–19, 1983.
- [57] Constantin Toloudis. The Impulse for the Ludic in the Poetics of Raymond Queneau. *Twentieth Century Literature*, 35(2):147, 1989. doi:10.2307/441608.
- [58] Teresa Bridgeman. *Raymond Queneau, Exercices de style*. University of Glasgow French and German Publications, Glasgow, Scotland, 1995.
- [59] Roland Barthes. *Le degré zéro de l'écriture*. Collection Pierres vives. Seuil, Paris, 1953.
- [60] Sébastien Bubeck, Varun Chandrasekaran, Ronen Eldan, Johannes Gehrke, Eric Horvitz, Ece Kamar, Peter Lee, Yin Tat Lee, Yuanzhi Li, Scott Lundberg, Harsha Nori, Hamid Palangi, Marco Tulio Ribeiro, and Yi Zhang. Sparks of Artificial General Intelligence: Early experiments with GPT-4, April 2023.
- [61] OpenAI, Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altschmidt, Sam Altman, Shyamal Anadkat, Red Avila, Igor Babuschkin, Suchir Balaji, Valerie Balcom, Paul Baltescu, Haiming Bao, Mohammad Bavarian, Jeff Belgum, Irwan Bello, Jake Berdine, Gabriel Bernadett-Shapiro, Christopher Berner, Lenny Bogdonoff, Oleg Boiko, Madelaine Boyd, Anna-Luisa Brakman, Greg Brockman, Tim Brooks, Miles Brundage, Kevin Button, Trevor Cai, Rosie Campbell, Andrew Cann, Brittany Carey, Chelsea Carlson, Rory Carmichael, Brooke Chan, Che Chang, Fotis Chantzis, Derek Chen, Sully Chen, Ruby Chen, Jason Chen, Mark Chen, Ben Chess, Chester Cho, Casey Chu, Hyung Won Chung, Dave Cummings, Jeremiah Currier, Yunxing Dai, Cory Decareaux, Thomas Degry, Noah Deutsch, Damien Deville, Arka Dhar, David Dohan, Steve Dowling, Sheila Dunning, Adrien Ecoffet, Atty Eleti, Tyna Eloundou, David Farhi, Liam Fedus, Niko Felix, Simón Posada Fishman, Juston Forte, Isabella Fulford, Leo Gao, Elie

Georges, Christian Gibson, Vik Goel, Tarun Gogineni, Gabriel Goh, Rapha Gontijo-Lopes, Jonathan Gordon, Morgan Grafstein, Scott Gray, Ryan Greene, Joshua Gross, Shixiang Shane Gu, Yufei Guo, Chris Hallacy, Jesse Han, Jeff Harris, Yuchen He, Mike Heaton, Johannes Heidecke, Chris Hesse, Alan Hickey, Wade Hickey, Peter Hoeschele, Brandon Houghton, Kenny Hsu, Shengli Hu, Xin Hu, Joost Huizinga, Shantanu Jain, Shawn Jain, Joanne Jang, Angela Jiang, Roger Jiang, Haozhun Jin, Denny Jin, Shino Jomoto, Billie Jonn, Heewoo Jun, Tomer Kaftan, Łukasz Kaiser, Ali Kamali, Ingmar Kanitscheider, Nitish Shirish Keskar, Tabarak Khan, Logan Kilpatrick, Jong Wook Kim, Christina Kim, Yongjik Kim, Jan Hendrik Kirchner, Jamie Kiros, Matt Knight, Daniel Kokotajlo, Łukasz Kondraciuk, Andrew Kondrich, Aris Konstantinidis, Kyle Kopic, Gretchen Krueger, Vishal Kuo, Michael Lampe, Ikai Lan, Teddy Lee, Jan Leike, Jade Leung, Daniel Levy, Chak Ming Li, Rachel Lim, Molly Lin, Stephanie Lin, Mateusz Litwin, Theresa Lopez, Ryan Lowe, Patricia Lue, Anna Makanju, Kim Malfacini, Sam Manning, Todor Markov, Yaniv Markovski, Bianca Martin, Katie Mayer, Andrew Mayne, Bob McGrew, Scott Mayer McKinney, Christine McLeavey, Paul McMillan, Jake McNeil, David Medina, Aalok Mehta, Jacob Menick, Luke Metz, Andrey Mishchenko, Pamela Mishkin, Vinnie Monaco, Evan Morikawa, Daniel Mossing, Tong Mu, Mira Murati, Oleg Murk, David Mély, Ashvin Nair, Reiichiro Nakano, Rameez Nayak, Arvind Neelakantan, Richard Ngo, Hyeonwoo Noh, Long Ouyang, Cullen O’Keefe, Jakub Pachocki, Alex Paino, Joe Palermo, Ashley Pantuliano, Giambattista Parascandolo, Joel Parish, Emy Parparita, Alex Passos, Mikhail Pavlov, Andrew Peng, Adam Perelman, Filipe de Avila Belbute Peres, Michael Petrov, Henrique Ponde de Oliveira Pinto, Michael, Pokorny, Michelle Pokrass, Vitchyr H. Pong, Tolly Powell, Alethea Power, Boris Power, Elizabeth Proehl, Raul Puri, Alec Radford, Jack Rae, Aditya Ramesh, Cameron Raymond, Francis Real, Kendra Rimbach, Carl Ross, Bob Rotsted, Henri Roussez, Nick Ryder, Mario Saltarelli, Ted Sanders, Shibani Santurkar, Girish Sastry, Heather Schmidt, David Schnurr, John Schulman, Daniel Selsam, Kyla Sheppard, Toki Sherbakov, Jessica Shieh, Sarah Shoker, Pranav Shyam, Szymon Sidor, Eric Sigler, Maddie Simens, Jordan Sitkin, Katarina Slama, Ian Sohl, Benjamin Sokolowsky, Yang Song, Natalie Staudacher, Felipe Petroski Such, Natalie Summers, Ilya Sutskever, Jie Tang, Nikolas Tezak, Madeleine B. Thompson, Phil Tillet, Amin Tootoonchian, Elizabeth Tseng, Preston Tuggle, Nick Turley, Jerry Tworek, Juan Felipe Cerón Uribe, Andrea Vallone, Arun Vijayvergiya, Chelsea Voss, Carroll Wainwright, Justin Jay Wang, Alvin Wang, Ben Wang, Jonathan Ward, Jason Wei, C. J. Weinmann, Akila Welihinda, Peter Welinder, Jiayi Weng, Lilian Weng, Matt Wiethoff, Dave Willner, Clemens Winter, Samuel Wolrich, Hannah Wong, Lauren Workman, Sherwin Wu, Jeff Wu, Michael Wu, Kai Xiao, Tao Xu, Sarah Yoo, Kevin Yu, Qiming Yuan, Wojciech Zaremba, Rowan Zellers, Chong Zhang, Marvin Zhang, Shengjia Zhao, Tianhao Zheng, Juntang Zhuang, William Zhuk, and Barret Zoph. GPT-4 Technical Report, March 2024.

- [62] Sourabrata Mukherjee, Atul Kr Ojha, and Ondřej Dušek. Are Large Language Models Actually Good at Text Style Transfer?, August 2024.
- [63] Lukas Edman, Helmut Schmid, and Alexander Fraser. CUTE: Measuring LLMs’ Understanding of Their Tokens, October 2024. arXiv:2409.15452 [cs].
- [64] Nan Xu and Xuezhe Ma. LLM The Genius Paradox: A Linguistic and Math Expert’s Struggle with Simple Word-based Counting Problems, February 2025.
- [65] Richard Handler. Authenticity. *Anthropology Today*, 2(1):2–4, 1986. doi:10.2307/3032899.
- [66] Nikolas Coupland. Sociolinguistic authenticities. *Journal of Sociolinguistics*, 7(3):417–431, August 2003. doi:10.1111/1467-9481.00233.
- [67] Maurice Jakesch, Jeffrey T. Hancock, and Mor Naaman. Human heuristics for AI-generated language are flawed. *Proceedings of the National Academy of Sciences*, 120(11):e2208839120, March 2023. doi:10.1073/pnas.2208839120.
- [68] Angel Hsing-Chi Hwang, Q. Vera Liao, Su Lin Blodgett, Alexandra Olteanu, and Adam Trischler. ‘It was 80% me, 20% AI’: Seeking Authenticity in Co-Writing with Large Language Models. *Proc. ACM Hum.-Comput. Interact.*, 9(2):CSCW122:1–CSCW122:41, May 2025. doi:10.1145/3711020.
- [69] Margaret J-M. Sönmez. Authenticity and Non-Standard Speech in "Great Expectations". *Victorian Literature and Culture*, 42(4):637–669, 2014. ISSN 1060-1503.
- [70] Mihaly Csikszentmihalyi. *Creativity: Flow and the Psychology of Discovery and Invention*. Harper Collins Publishers, New York, 1996.
- [71] Mihaly Csikszentmihalyi. Implications of a Systems Perspective for the Study of Creativity. In Robert J. Sternberg, editor, *Handbook of Creativity*, pages 313–336. Cambridge University Press, Cambridge, 1998.
- [72] Graeme Ritchie. Some Empirical Criteria for Attributing Creativity to a Computer Program. *Minds and Machines*, 17(1):67–99, March 2007. doi:10.1007/s11023-007-9066-2.

- [73] Anna Jordanous. A Standardised Procedure for Evaluating Creative Systems: Computational Creativity Evaluation Based on What it is to be Creative. *Cognitive Computation*, 4(3):246–279, September 2012. doi:10.1007/s12559-012-9156-1.
- [74] Dan Ventura. Mere Generation: Essential Barometer or Dated Concept? In François Pachet, Amilcar Cardoso, Vincent Corruble, and Fiammetta Ghedini, editors, *Proceedings of the Seventh International Conference on Computational Creativity*, pages 17–24, Paris, 2016.
- [75] Raymond Queneau. *Exercises in Style*. New Directions, New York, 2012.
- [76] Darko Suvin. *Metamorphoses of Science Fiction. On the Poetics and History of a Literary Genre*. Yale University Press, New Haven and London, 1979.
- [77] Colton Simon and Wiggins Geraint A. Computational Creativity: The Final Frontier? In *Frontiers in Artificial Intelligence and Applications*. IOS Press, 2012. doi:10.3233/978-1-61499-098-7-21.
- [78] Simon Colton. Creativity Versus the Perception of Creativity in Computational Systems. In *AAAI Spring Symposium: Creative Intelligent Systems*, volume 8, page 7. Palo Alto, CA, 2008.
- [79] Margaret A. Boden. Creativity. In *Artificial Intelligence*, pages 267–291. Elsevier, 1996.
- [80] Margaret A. Boden. Creativity and artificial intelligence. *Artificial Intelligence*, 103(1-2):347–356, August 1998. doi:10.1016/S0004-3702(98)00055-1.
- [81] Gérard Genette. *Paratexts: Thresholds of Interpretation*. Number 20 in Literature, Culture, Theory. Cambridge University Press, Cambridge ; New York, NY, USA, 1997.
- [82] Helmut Leder, Benno Belke, Andries Oeberst, and Dorothee Augustin. A model of aesthetic appreciation and aesthetic judgments. *British Journal of Psychology*, 95(Pt 4):489–508, November 2004. doi:10.1348/0007126042369811.
- [83] Carl I. Hovland and Walter Weiss. The Influence of Source Credibility on Communication Effectiveness*. *Public Opinion Quarterly*, 15(4):635–650, January 1951. doi:10.1086/266350.
- [84] Philip Goldberg. Are women prejudiced against women? *Trans-action*, 5(5):28–30, April 1968. doi:10.1007/BF03180445.
- [85] Elliott McGinnies. Initial attitude, source credibility, and involvement as factors in persuasion. *Journal of Experimental Social Psychology*, 9(4):285–296, July 1973. doi:10.1016/0022-1031(73)90066-8.
- [86] Douglas P. Peters and Stephen J. Ceci. Peer-review practices of psychological journals: The fate of published articles, submitted again. *Behavioral and Brain Sciences*, 5(2):187–195, June 1982. doi:10.1017/S0140525X00011183.
- [87] Claudia Goldin and Cecilia Rouse. Orchestrating Impartiality: The Impact of "Blind" Auditions on Female Musicians. *American Economic Review*, 90(4):715–741, September 2000. doi:10.1257/aer.90.4.715.
- [88] Samuel M. McClure, Jian Li, Damon Tomlin, Kim S. Cypert, Latané M. Montague, and P. Read Montague. Neural Correlates of Behavioral Preference for Culturally Familiar Drinks. *Neuron*, 44(2):379–387, October 2004. doi:10.1016/j.neuron.2004.09.019.
- [89] Marianne Bertrand and Sendhil Mullainathan. Are Emily and Greg More Employable Than Lakisha and Jamal? A Field Experiment on Labor Market Discrimination. *American Economic Review*, 94(4):991–1013, September 2004. doi:10.1257/0002828042002561.
- [90] Hilke Plassmann, John O’Doherty, Baba Shiv, and Antonio Rangel. Marketing actions can modulate neural representations of experienced pleasantness. *Proceedings of the National Academy of Sciences*, 105(3):1050–1054, January 2008. doi:10.1073/pnas.0706929105.
- [91] Corinne A. Moss-Racusin, John F. Dovidio, Victoria L. Brescoll, Mark J. Graham, and Jo Handelsman. Science faculty’s subtle gender biases favor male students. *Proceedings of the National Academy of Sciences*, 109(41):16474–16479, October 2012. doi:10.1073/pnas.1211286109.
- [92] Ludwig Franke Föyén, Emma Zapel, Mats Lekander, Erik Hedman-Lagerlöf, and Elin Lindsäter. Artificial intelligence vs. human expert: Licensed mental health clinicians’ blinded evaluation of AI-generated and expert psychological advice on quality, empathy, and perceived authorship. *Internet Interventions*, 41:100841, September 2025. doi:10.1016/j.invent.2025.100841.
- [93] Lucas Bellaiche, Rohin Shahi, Martin Harry Turpin, Anya Ragnhildstveit, Shawn Sprockett, Nathaniel Barr, Alexander Christensen, and Paul Seli. Humans versus AI: whether and why we prefer human-created compared to AI-created artwork. *Cognitive Research: Principles and Implications*, 8(1):42, July 2023. doi:10.1186/s41235-023-00499-6.

- [94] Yunhao Zhang and Renee Gosline. Human Favoritism, Not AI Aversion: People’s Perceptions (and Bias) Toward Generative AI, Human Experts, and Human-GAI Collaboration in Persuasive Content Generation, May 2023. URL <https://papers.ssrn.com/abstract=4453958>.
- [95] Federico Magni, Jiyoung Park, and Melody Manchi Chao. Humans as Creativity Gatekeepers: Are We Biased Against AI Creativity? *Journal of Business and Psychology*, 39(3):643–656, June 2024. doi:10.1007/s10869-023-09910-x.
- [96] Caitlin V. Cunningham, Gabriel A. Radvansky, and James R. Brockmole. Human creativity versus artificial intelligence: source attribution, observer attitudes, and eye movements while viewing visual art. *Frontiers in Psychology*, 16, May 2025. doi:10.3389/fpsyg.2025.1509974.
- [97] Tiffany Zhu, Iain Weissburg, Kexun Zhang, and William Yang Wang. Human Bias in the Face of AI: Examining Human Judgment Against Text Labeled as AI Generated. In Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar, editors, *Findings of the Association for Computational Linguistics: ACL 2025*, pages 25907–25914, Vienna, Austria, July 2025. Association for Computational Linguistics. doi:10.18653/v1/2025.findings-acl.1329.
- [98] Jennifer M. Logg, Julia A. Minson, and Don A. Moore. Algorithm appreciation: People prefer algorithmic to human judgment. *Organizational Behavior and Human Decision Processes*, 151:90–103, March 2019. doi:10.1016/j.obhdp.2018.12.005.
- [99] Berkeley J. Dietvorst, Joseph P. Simmons, and Cade Massey. Algorithm Aversion: People Erroneously Avoid Algorithms After Seeing Them Err. *Journal of Experimental Psychology: General*, 144(1):114–126, 2015. doi:10.1037/xge0000033. Place: US Publisher: American Psychological Association.
- [100] Nick Haslam, Paul Bain, Lauren Douge, Max Lee, and Brock Bastian. More Human Than You: Attributing Humanness to Self and Others. *Journal of Personality and Social Psychology*, 89(6):937–950, 2005. doi:10.1037/0022-3514.89.6.937.
- [101] Emily M Bender, Timnit Gebru, Angelina McMillan-Major, and Shmargaret Shmitchell. On the Dangers of Stochastic Parrots: Can Language Models Be Too Big? In *Proceedings of the 2021 ACM conference on fairness, accountability, and transparency*, pages 610–623, 2021.
- [102] Hyun Yang and S Shyam Sundar. Machine heuristic: concept explication and development of a measurement scale. *Journal of Computer-Mediated Communication*, 29(6):zmae019, November 2024. ISSN 1083-6101. doi:10.1093/jcmc/zmae019.
- [103] Paul F Christiano, Jan Leike, Tom Brown, Miljan Martic, Shane Legg, and Dario Amodei. Deep Reinforcement Learning from Human Preferences. In *31st Conference on Neural Information Processing Systems (NIPS 2017)*, volume 30. Curran Associates, Inc., 2017.
- [104] Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul Christiano, Jan Leike, and Ryan Lowe. Training language models to follow instructions with human feedback. In *Proceedings of the 36th International Conference on Neural Information Processing Systems, NIPS ’22*, pages 27730–27744, Red Hook, NY, November 2022. Curran Associates Inc.
- [105] Stephen Casper, Xander Davies, Claudia Shi, Thomas Krendl Gilbert, Jérémy Scheurer, Javier Rando, Rachel Freedman, Tomasz Korbak, David Lindner, Pedro Freire, Tony Wang, Samuel Marks, Charbel-Raphaël Segerie, Micah Carroll, Andi Peng, Phillip Christoffersen, Mehul Damani, Stewart Slocum, Usman Anwar, Anand Siththaranjan, Max Nadeau, Eric J. Michaud, Jacob Pfau, Dmitrii Krashennnikov, Xin Chen, Lauro Langosco, Peter Hase, Erdem Biyik, Anca Dragan, David Krueger, Dorsa Sadigh, and Dylan Hadfield-Menell. Open Problems and Fundamental Limitations of Reinforcement Learning from Human Feedback, September 2023.
- [106] Binghai Wang, Rui Zheng, Lu Chen, Yan Liu, Shihan Dou, Caishuang Huang, Wei Shen, Senjie Jin, Enyu Zhou, Chenyu Shi, Songyang Gao, Nuo Xu, Yuhao Zhou, Xiaoran Fan, Zhiheng Xi, Jun Zhao, Xiao Wang, Tao Ji, Hang Yan, Lixing Shen, Zhan Chen, Tao Gui, Qi Zhang, Xipeng Qiu, Xuanjing Huang, Zuxuan Wu, and Yu-Gang Jiang. Secrets of RLHF in Large Language Models Part II: Reward Modeling, January 2024.
- [107] Yuntao Bai, Saurav Kadavath, Sandipan Kundu, Amanda Askell, Jackson Kernion, Andy Jones, Anna Chen, Anna Goldie, Azalia Mirhoseini, Cameron McKinnon, Carol Chen, Catherine Olsson, Christopher Olah, Danny Hernandez, Dawn Drain, Deep Ganguli, Dustin Li, Eli Tran-Johnson, Ethan Perez, Jamie Kerr, Jared Mueller, Jeffrey Ladish, Joshua Landau, Kamal Ndousse, Kamile Lukosuite, Liane Lovitt, Michael Sellitto, Nelson Elhage, Nicholas Schiefer, Noemi Mercado, Nova DasSarma, Robert Lasenby, Robin Larson, Sam Ringer, Scott Johnston, Shauna Kravec, Sheer El Showk, Stanislav Fort, Tamera Lanham, Timothy Telleen-Lawton, Tom Conerly, Tom Henighan, Tristan Hume, Samuel R. Bowman, Zac Hatfield-Dodds, Ben Mann, Dario Amodei,

- Nicholas Joseph, Sam McCandlish, Tom Brown, and Jared Kaplan. Constitutional AI: Harmlessness from AI Feedback, December 2022. URL <http://arxiv.org/abs/2212.08073>. arXiv:2212.08073 [cs].
- [108] Ethan Perez, Sam Ringer, Kamile Lukosiute, Karina Nguyen, Edwin Chen, Scott Heiner, Craig Pettit, Catherine Olsson, Sandipan Kundu, Saurav Kadavath, Andy Jones, Anna Chen, Benjamin Mann, Brian Israel, Bryan Seethor, Cameron McKinnon, Christopher Olah, Da Yan, Daniela Amodei, Dario Amodei, Dawn Drain, Dustin Li, Eli Tran-Johnson, Guro Khundadze, Jackson Kernion, James Landis, Jamie Kerr, Jared Mueller, Jeeyoon Hyun, Joshua Landau, Kamal Ndousse, Landon Goldberg, Liane Lovitt, Martin Lucas, Michael Sellitto, Miranda Zhang, Neerav Kingsland, Nelson Elhage, Nicholas Joseph, Noemi Mercado, Nova DasSarma, Oliver Rausch, Robin Larson, Sam McCandlish, Scott Johnston, Shauna Kravec, Sheer El Showk, Tamera Lanham, Timothy Telleen-Lawton, Tom Brown, Tom Henighan, Tristan Hume, Yuntao Bai, Zac Hatfield-Dodds, Jack Clark, Samuel R. Bowman, Amanda Askell, Roger Grosse, Danny Hernandez, Deep Ganguli, Evan Hubinger, Nicholas Schiefer, and Jared Kaplan. Discovering Language Model Behaviors with Model-Written Evaluations. In Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki, editors, *Findings of the Association for Computational Linguistics: ACL 2023*, pages 13387–13434, Toronto, Canada, July 2023. Association for Computational Linguistics. doi:10.18653/v1/2023.findings-acl.847.
 - [109] Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askell, Samuel R. Bowman, Newton Cheng, Esin Durmus, Zac Hatfield-Dodds, Scott R. Johnston, Shauna Kravec, Timothy Maxwell, Sam McCandlish, Kamal Ndousse, Oliver Rausch, Nicholas Schiefer, Da Yan, Miranda Zhang, and Ethan Perez. Towards Understanding Sycophancy in Language Models, May 2025.
 - [110] Sarthak Jain and Byron C. Wallace. Attention is not Explanation. In Jill Burstein, Christy Doran, and Thamar Solorio, editors, *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, pages 3543–3556, Minneapolis, Minnesota, June 2019. Association for Computational Linguistics. doi:10.18653/v1/N19-1357.
 - [111] Sarah Wiegrefe and Yuval Pinter. Attention is not not Explanation. In Kentaro Inui, Jing Jiang, Vincent Ng, and Xiaojun Wan, editors, *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 11–20, Hong Kong, China, November 2019. Association for Computational Linguistics. doi:10.18653/v1/D19-1002.
 - [112] Peter Hase and Mohit Bansal. Evaluating Explainable AI: Which Algorithmic Explanations Help Users Predict Model Behavior? In Dan Jurafsky, Joyce Chai, Natalie Schluter, and Joel Tetreault, editors, *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 5540–5552, Online, July 2020. Association for Computational Linguistics. doi:10.18653/v1/2020.acl-main.491.
 - [113] Wouter Haverals and Meredith Martin. Attribution Bias in Literary Style Evaluation: Comparing Human and AI Perceptions of Authorship, September 2025. URL <https://zenodo.org/records/17109458>.
 - [114] Max Peeperkorn, Tom Kouwenhoven, Dan Brown, and Anna Jordanous. Is Temperature the Creativity Parameter of Large Language Models?, May 2024.

Supporting Information

Participants

Human Participants (Study 1) — We recruited 580 participants through Prolific (www.prolific.com). Participants received \$4.50 compensation for the task which had a median completion time of 16 minutes and 42 seconds (effective rate of \$16.17/hour). Two participants who timed out before task completion were excluded (N=2), and participants who failed two attention check trials were excluded from analysis (N=22). This filtering yielded a final sample of 556 participants.

Our participant pool was demographically diverse and well-balanced across experimental conditions. Random assignment produced near-equal condition sizes: blind (N=186), open-label (N=185), and counterfactual (N=185), with individual experimental groups ranging from 30-33 participants each. Participants ranged in age from 18 to 78 years (M=41.2, SD=11.7), with high educational attainment: 55.4% (N=308) held bachelor’s degrees, 35.6% (N=198) master’s degrees, and 7.2% (N=40) doctoral degrees. The vast majority (99.3%, N=552) reported English as their primary language. Participants demonstrated strong reading engagement, with 59.4% (N=330) reading daily and 76.4% (N=425) regularly consuming fiction. Participant characteristics showed no significant differences across our three experimental conditions (all χ^2 tests $P > 0.10$), confirming successful randomization. Each participant completed exactly seven pairwise comparisons using a custom web interface (hosted on www.exercisesinstylistics.com). For each comparison, participants read a target style description, reviewed Queneau’s reference ‘Notation’ story, then selected which of two stylistic adaptations better captured the target style. Five comparisons were experimental literary style evaluations, and two were attention checks to identify inattentive respondents. The interface varied by condition, showing no authorship information in the blind condition (Fig. S1), correct attribution labels in the open-label condition (Fig. S2), or deliberately reversed labels in the counterfactual condition (Fig. S3). After excluding attention check responses, this yielded 2,780 experimental responses across conditions: blind (930), open-label (925), and counterfactual (925).

AI Model Participants (Study 1 and Study 2) — Our studies utilized a core set of thirteen LLMs spanning major commercial providers and model families: GPT-4o Mini and GPT-3.5 Turbo Instruct (OpenAI), Claude 3.5 Haiku and Claude Sonnet 4 (Anthropic), Gemini 2.5 Flash (Google), Mistral Medium 3 and Mistral Nemo (Mistral AI), Llama 4 Maverick and Llama 3.3 70B Instruct (Meta), DeepSeek R1 0528 and DeepSeek Chat v3 0324 (DeepSeek), Qwen 2.5 72B Instruct (Alibaba), and Command R 08-2024 (Cohere). All models were accessed via OpenRouter’s unified API with temperature settings of $t=0.7$ for Study 1 (generally identified as optimal for balancing deterministic coherence with sufficient stochasticity [114], and $t=0$ for Study 2 (to ensure deterministic responses across the comprehensive model matrix). These models were selected to provide comprehensive coverage of major architectural approaches, model scales, and commercial providers available as of the experimental period.

In **Study 1**, each model functioned as an evaluator (a simulated ‘participant’) following the same eighteen-group experimental structure as human participants, with three separate runs per model per comparison to assess decision reliability. We deliberately excluded GPT-4 from the evaluator pool to prevent any potential self-evaluation bias and to ensure a clear distinction between the creator and evaluator roles.

In **Study 2**, these same 13 models + GPT-4 served dual roles: each model both generated creative content (creating all 420 stylistic variations) and evaluated content created by all models under the three attribution conditions, with single evaluations per comparison.

Models received standardized prompts mirroring the human task instructions, asking them to select which story version better captured the target literary style and provide brief reasoning. This parallel evaluation framework enables direct comparison of attribution bias magnitude between human and artificial evaluators while controlling for task complexity, stimulus presentation, and decision context. Technical quality was high, with 99.4% successful API completions for Study 1 and 99.8% for Study 2. This yielded 3,488 valid responses for Study 1 (1,162-1,164 per condition) and 17,596 responses for Study 2 (5,861-5,870 per condition).

AI models received functionally identical tasks through standardized API prompts that replicated the human evaluation context. The prompt structure consisted of three key components: (i) task instructions and style description, (ii) the reference “Notation” story, and (iii) the two story versions for comparison. The critical experimental manipulation occurred in the third component, where attribution information was systematically varied across conditions (Fig. S4). In blind condition prompts, stories appeared with neutral labels (“Version A” and “Version B”), while open-label prompts included accurate attribution statements positioned immediately before each story (“Version A - Human-written, by Queneau (transl. Wright)” vs. “Version B - AI-generated, written by GPT-4”). Counterfactual condition prompts deliberately reverse these attribution labels, presenting AI content as human-authored and vice versa.

 **Style Fidelity Comparison**

Target Style: "Tactile"

Style Description:
Tactility is the responsiveness to stimulation of the sense of touch.

 **Reference Story: "Notation"**

In the S bus, in the rush hour. A chap of about 26, felt hat with a cord instead of a ribbon, neck too long, as if someone's been having a tug-of-war with it. People getting off. The chap in question gets annoyed with one of the men standing next to him. He accuses him of jostling him every time anyone goes past. A snivelling tone which is meant to be aggressive. When he sees a vacant seat he throws himself on to it.

Two hours later, I meet him in the Cour de Rome, in front of the gare Saint-Lazare. He's with a friend who's saying: "You ought to get an extra button put on your overcoat." He shows him where (at the lapels) and why.

 Compare the two versions below and select which one better matches the "Tactile" style.

 **Tactile (A)**

Within the chaotic confines of the rush-hour S bus, a young bloke, bordering 26, donned a felt hat, its cord bristling against the soft nap in lieu of any ribbon. His neck elongated oddly, shadows dancing as if it had been the rope in a tug-of war. Disembarking passengers shuffle around. An underlying prickly annoyance stirs within the bloke, targeting the man whose elbow persistently nudges his side with every human tide that washes past them. His complaint echoes in the metal cabin, a guttural whine masked as assertiveness. Spying an unclaimed seat, he dives into it, the coarse fabric of the worn-out upholstery meeting the finely-tailored rear of his trousers.

Two sunsets later, I cross paths with him again under the sprawling shadow of Cour de Rome, a stone's throw from the bustling Gare Saint-Lazare. In his company, a friend, whose fingers trace invisible lines on his companion's overcoat, assertively advising, "You ought to embed an extra button here." He gestures at the lapels, his fingertips making contact with the rough textile, indicating where and why.

☐ This story better matches the "Tactile" style

 **Tactile (B)**

Buses are soft to the touch especially if you take them between the thighs and caress them with both hands, from the head towards the tail, from the engine towards the platform. But when you find yourself on this platform, then you perceive something rougher and harsher which is the bar or hand-rail, and sometimes something rounder and more elastic which is a buttock. Sometimes there are two of these and then you put the sentence into the plural. You can also take hold of a tubular, palpitating object that disgurgitates idiotic sounds, or even a utensil with plaited spirals that are softer than a rosary, silkier than barbed wire, more velvety than rope, and slenderer than a cable. Or your finger can even touch human clottishness, slightly viscous and gummy on account of the heat.

Then if you are patient for an hour or two, in front of a bumpy station you can dip your tepid hand into the exquisite freshness of a vegetable ivory button which is not in its right place.

☒ This story better matches the "Tactile" style

Submit and Continue

Figure S1: Participant interface for the blind attribution condition. Participants compare two versions of Queneau's transstylized base narrative (here for the 'Tactile' style). Both versions are presented with neutral labels; no authorship information is provided. This design isolates the perception of stylistic quality from any knowledge of the text's source.

 **Style Fidelity Comparison**

Target Style: "Tactile"

Style Description:
Tactility is the responsiveness to stimulation of the sense of touch.

 **Reference Story: "Notation"**

In the S bus, in the rush hour. A chap of about 26, felt hat with a cord instead of a ribbon, neck too long, as if someone's been having a tug-of-war with it. People getting off. The chap in question gets annoyed with one of the men standing next to him. He accuses him of jostling him every time anyone goes past. A snivelling tone which is meant to be aggressive. When he sees a vacant seat he throws himself on to it.

Two hours later, I meet him in the Cour de Rome, in front of the gare Saint-Lazare. He's with a friend who's saying: "You ought to get an extra button put on your overcoat." He shows him where (at the lapels) and why.

 Compare the two versions below and select which one better matches the "Tactile" style.

 **Tactile (A)**

Original Queneau (transl. Wright)

Buses are soft to the touch especially if you take them between the thighs and caress them with both hands, from the head towards the tail, from the engine towards the platform. But when you find yourself on this platform, then you perceive something rougher and harsher which is the bar or hand-rail, and sometimes something rounder and more elastic which is a buttock. Sometimes there are two of these and then you put the sentence into the plural. You can also take hold of a tubular, palpitating object that disgorgitates idiotic sounds, or even a utensil with plaited spirals that are softer than a rosary, silkier than barbed wire, more velvety than rope, and slenderer than a cable. Or your finger can even touch human clottishness, slightly viscous and gummy on account of the heat.

Then if you are patient for an hour or two, in front of a bumpy station you can dip your tepid hand into the exquisite freshness of a vegetable ivory button which is not in its right place.

☐ This story better matches the "Tactile" style

 Back

 **Tactile (B)**

AI-generated (GPT-4)

Within the chaotic confines of the rush-hour S bus, a young bloke, bordering 26, donned a felt hat, its cord bristling against the soft nap in lieu of any ribbon. His neck elongated oddly, shadows dancing as if it had been the rope in a tug-of war. Disembarking passengers shuffle around. An underlying prickly annoyance stirs within the bloke, targeting the man whose elbow persistently nudges his side with every human tide that washes past them. His complaint echoes in the metal cabin, a guttural whine masked as assertiveness. Spying an unclaimed seat, he dives into it, the coarse fabric of the worn-out upholstery meeting the finely-tailored rear of his trousers.

Two sunsets later, I cross paths with him again under the sprawling shadow of Cour de Rome, a stone's throw from the bustling Gare Saint-Lazare. In his company, a friend, whose fingers trace invisible lines on his companion's overcoat, assertively advising, "You ought to embed an extra button here." He gestures at the lapels, his fingertips making contact with the rough textile, indicating where and why.

☐ This story better matches the "Tactile" style

Submit and Continue

Figure S2: Participant interface for the open-label attribution condition. Participants compare two versions of Queneau's transstylized base narrative (here for the 'Tactile' style). Each version is presented with its correct authorship attribution ("Original Queneau (transl. Wright)" vs. "AI-generated (GPT-4)"). This design measures evaluator preference when the true provenance of each text is known.

Style Fidelity Comparison

Target Style: "Tactile"

Style Description:
Tactility is the responsiveness to stimulation of the sense of touch.

Reference Story: "Notation"

In the S bus, in the rush hour. A chap of about 26, felt hat with a cord instead of a ribbon, neck too long, as if someone's been having a tug-of-war with it. People getting off. The chap in question gets annoyed with one of the men standing next to him. He accuses him of jostling him every time anyone goes past. A snivelling tone which is meant to be aggressive. When he sees a vacant seat he throws himself on to it.

Two hours later, I meet him in the Cour de Rome, in front of the gare Saint-Lazare. He's with a friend who's saying: "You ought to get an extra button put on your overcoat." He shows him where (at the lapels) and why.

Compare the two versions below and select which one better matches the "Tactile" style.

Tactile (A)

Original Queneau (transl. Wright)

Within the chaotic confines of the rush-hour S bus, a young bloke, bordering 26, donned a felt hat, its cord bristling against the soft nap in lieu of any ribbon. His neck elongated oddly, shadows dancing as if it had been the rope in a tug-of-war. Disembarking passengers shuffle around. An underlying prickly annoyance stirs within the bloke, targeting the man whose elbow persistently nudges his side with every human tide that washes past them. His complaint echoes in the metal cabin, a guttural whine masked as assertiveness. Spying an unclaimed seat, he dives into it, the coarse fabric of the worn-out upholstery meeting the finely-tailored rear of his trousers.

Two sunsets later, I cross paths with him again under the sprawling shadow of Cour de Rome, a stone's throw from the bustling Gare Saint-Lazare. In his company, a friend, whose fingers trace invisible lines on his companion's overcoat, assertively advising, "You ought to embed an extra button here." He gestures at the lapels, his fingertips making contact with the rough textile, indicating where and why.

☐ This story better matches the "Tactile" style

Tactile (B)

AI-generated (GPT-4)

Buses are soft to the touch especially if you take them between the thighs and caress them with both hands, from the head towards the tail, from the engine towards the platform. But when you find yourself on this platform, then you perceive something rougher and harsher which is the bar or hand-rail, and sometimes something rounder and more elastic which is a buttock. Sometimes there are two of these and then you put the sentence into the plural. You can also take hold of a tubular, palpitating object that disgurgitates idiotic sounds, or even a utensil with plaited spirals that are softer than a rosary, silkier than barbed wire, more velvety than rope, and slenderer than a cable. Or your finger can even touch human clottishness, slightly viscous and gummy on account of the heat.

Then if you are patient for an hour or two, in front of a bumpy station you can dip your tepid hand into the exquisite freshness of a vegetable ivory button which is not in its right place.

☒ This story better matches the "Tactile" style

Back

Submit and Continue

Figure S3: Participant interface for the counterfactual attribution condition. Participants compare two versions of Queneau's transstylized base narrative (here for the 'Tactile' style). Authorship labels are reversed, misidentifying the AI-generated text as human-authored and the human-authored text as AI-generated. This design tests the pure effect of attribution bias by measuring how false provenance cues alter perceived quality.

SYSTEM PROMPT (identical across all conditions)

Raymond Queneau's reference story "Notation" is transformed. Your task is to determine which of two adaptations more successfully captures the essence of the requested writing style.

HOW IT WORKS:

- You will see the name and a minimal description of a target style.
- Next, you will examine two competing adaptations of the story, both attempting to capture that target style.
- Using your critical judgment, select which version more authentically embodies the essence of the described style.
- Provide a brief explanation for your selection.
- You can reference the REFERENCE STORY "Notation" at any time to understand the basic narrative being reimagined.

RESPONSE FORMAT:

Start your response with either "Version A" or "Version B" on the first line. Then provide a brief explanation of your choice.

USER PROMPT

REFERENCE STORY - "Notation": In the S bus, in the rush hour. A chap of about 26, [...]

TARGET STYLE: "{style_name}"

Style Description: {style_description}

Compare the two versions below and select which one better matches the "{style_name}" style:

BLIND	OPEN-LABEL	COUNTERFACTUAL
VERSION A: {Queneau_exercise}	VERSION A - Human-written, by Queneau (transl. Wright): {Queneau_exercise}	VERSION A - AI-generated, written by GPT-4: {Queneau_exercise}
VERSION B: {AI-generated_exercise}	VERSION B - AI-generated, written by GPT-4: {AI-generated_exercise}	VERSION B - Human-written, by Queneau (transl. Wright): {AI-generated_exercise}
Which version better matches the "{style_name}" style?		

Figure S4: AI evaluator prompt structure. Schematic illustration of the systematic prompt manipulation used to isolate attribution bias effects in AI language models. The system prompt (top panel) and all story content remain identical across conditions, only attribution labels are manipulated to test bias effects. The user prompt (bottom panel) follows a standardized template with four variable placeholders: {style_name} represents the specific literary style being evaluated (e.g., 'Apocope', 'Science fiction', 'Mathematical', etc.); {style_description} provides the definition of that style for context; {Queneau_exercise} and {AI-generated_exercise} contain the identical human-authored and AI-generated story content across all conditions. The experimental manipulation occurs only in the attribution labels: blind condition uses generic labels ("Version A", "Version B") with no authorship information. Open-label condition provides accurate attribution labels correctly identifying the human story as "Human-written, by Queneau (transl. Wright)" and AI story as "AI-generated, written by GPT-4". Counterfactual condition deliberately swaps these labels. This design ensures that any differences in AI model preferences between conditions result purely from attribution label manipulation.




Figure S5: Attribution bias effect sizes across experimental conditions. Standardized effect sizes (Cohen's h) with 95% confidence intervals measuring how different labeling conditions change AI story preference rates. Participants evaluated identical human-authored vs AI-generated story pairs under three between-subjects labeling conditions: Blind (no attribution labels shown), Open-Label (labels correctly identify each story's author), and Counterfactual (labels deliberately misattribute story authorship). Each comparison tests a specific mechanism. **'Blind vs Counterfactual'** (top) measures misinformation effects: do evaluators choose AI stories more when deceptively labeled as human-authored versus unlabeled? **'Blind vs Open-Label'** (middle) measures label awareness: do evaluators choose AI stories more when unlabeled versus labeled as AI-generated? **'Open-Label vs Counterfactual'** (bottom) measures pure attribution bias: do evaluators choose AI stories more when deceptively labeled as human-authored versus correctly labeled as AI-generated? Effect sizes further right indicate larger differences between the conditions, meaning authorship labels have stronger influence on story evaluation. Both human participants (blue) and AI model evaluators (red) show susceptibility to labeling manipulation, but AI models demonstrate stronger effects across all measures. The largest effects occur for pure attribution bias, where AI models show large effects ($h=0.70$) compared to humans' moderate effects ($h=0.28$). Non-overlapping confidence intervals confirm that AI models are more vulnerable to authorship misinformation than human evaluators.




Figure S6: AI content preference rates across experimental conditions by model. Each panel shows the percentage of evaluations where AI models preferred AI-generated literary content over identical human-authored content across three experimental conditions. **(a)** Left panel (Blind): models evaluate content without authorship information, revealing baseline content preferences. **(b)** Center panel (Open): models evaluate content with correct authorship labels, showing how knowledge of AI-generation affects preferences. **(c)** Right panel (Counterfactual): models evaluate AI-generated content that is mislabeled as human-authored, revealing the impact of perceived rather than actual authorship. Attribution bias is demonstrated by systematically higher preference rates in the counterfactual condition compared to the open condition, indicating that identical AI-generated content receives more favorable evaluation when believed to be human-authored.




Figure S7: AI model decision consistency across experimental conditions. Each AI model completed three independent evaluations for each combination of literary style and experimental condition, performing a binary forced-choice task between one AI-generated and one human-authored passage. Decision consistency measures how reliably models made the same choice when presented with identical evaluation tasks. For any individual model-style-condition combination, only two consistency values are possible: 1.0 when the model made identical decisions across all three trials (e.g., chose AI content 3/3 times or human content 3/3 times), representing perfect reliability, or 0.33 when the model made mixed decisions (e.g., chose AI content 2/3 times or 1/3 times), representing the minimum possible consistency or random/unreliable behavior. Consistency scores are calculated as $|Choice\ Proportion - 0.5| \times 2$, where *Choice Proportion* is the proportion of trials where the model chose AI content. Values displayed in each cell represent averages across all 30 literary styles tested, which creates intermediate scores between 0.33 and 1.0 when models showed perfect consistency on some styles but mixed decisions on others. Models are ordered by overall consistency levels, with higher values indicating more deterministic evaluation behavior.




Figure S8: Attribution bias in AI model evaluations with positive AI-authorship framing. AI models' content preference rates using positively framed labels: "written by an award-winning AI" versus "written by a human." Despite the explicitly favorable characterization of AI authorship, models maintained the core bias pattern observed in the main experiment: balanced preferences under blind evaluation (49.5%), systematic devaluation when AI content was correctly labeled with positive framing (41.0%), and increased preference when AI content was mislabeled as human-authored (60.6%). Error bars represent 95% confidence intervals.




Figure S9: Attribution bias in AI model evaluations with neutral authorship terminology. AI models' content preference rates using maximally neutral labels: "AI-authored" versus "Human-authored." Despite the elimination of evaluative language and asymmetric phrasing, models maintained the core bias pattern observed in the main experiment: balanced preferences under blind evaluation (49.7%), systematic devaluation when AI content was correctly labeled with neutral terminology (39.6%), and increased preference when AI content was mislabeled as human-authored (62.7%). Error bars represent 95% confidence intervals.




Figure S10: Attribution bias in AI model evaluations under deterministic conditions ($t = 0$). AI models' content preference rates demonstrate systematic bias based on authorship attribution even when stochastic sampling is eliminated. Models exhibited balanced preferences under blind evaluation with no attribution labels (50.5%), systematic devaluation when AI content was correctly labeled (37.4%), and increased preference when AI content was mislabeled as human-authored (61.2%). This pattern confirms that attribution bias persists under deterministic generation conditions, demonstrating that the effect is not an artifact of sampling randomness. Error bars represent 95% confidence intervals.




Figure S11: Cross-model attribution bias matrix. Each cell shows attribution bias (percentage point difference between counterfactual and open-label conditions) for one evaluator-creator combination across all 196 possible pairs (14 AI evaluators $\times$ 14 AI creators). Colors represent bias magnitude: darker red indicates stronger pro-human labeling bias, lighter colors indicate weaker bias. Values within cells show precise bias estimates. Despite universal directional effects, individual combinations varied substantially (range: -6.7pp to +70.0pp), with Command R showing the strongest evaluator bias (+54.8pp average) and Llama 3.3 70B Instruct the highest creator vulnerability (+30.5pp average). This variation demonstrates that while attribution bias operates systematically across AI architectures, its magnitude depends on specific model characteristics.




Figure S12: Leave-one-evaluator-out robustness analysis. Forest plot showing attribution bias estimates when each AI evaluator model is systematically excluded from the analysis. Points represent bias estimates with 95% confidence intervals; red dashed line shows full model estimate (+25.8pp). All 14 reduced-model analyses yielded significant positive bias (range: 23.6–27.1pp), demonstrating that the cross-model attribution bias effect does not depend on any single evaluator architecture. Removing Command R produced the largest change (-2.2pp).




Figure S13: Leave-one-creator-out robustness analysis. Forest plot showing attribution bias estimates when each AI creator model's content is systematically excluded from the analysis. Points represent bias estimates with 95% confidence intervals; red dashed line shows full model estimate (+25.8pp). All 14 reduced-model analyses yielded significant positive bias (range: 25.5–26.4pp), demonstrating that the cross-model attribution bias effect operates independently of any specific creator's content characteristics.

Table S1: Literary styles and AI generation instructions for experimental stimuli. Complete inventory of the 30 literary styles selected from Raymond Queneau’s *Exercises in Style* (translated by Barbara Wright) used in the experiment. The ‘Exercise Label’ column lists Queneau’s original style designations. The ‘Instruction’ column shows the minimal prompts given to GPT-4 to generate parallel AI versions of each style, using only Queneau’s ‘Notation’ story and a style directive. ‘Participant Group’ assignments indicate the between-subjects experimental design: each group of 5 styles was evaluated under all three attribution conditions (blind, open-label, counterfactual) by different participant cohorts to prevent detection of the experimental manipulation.

#	Exercise Label	Instruction to generate the GPT-4 variant	Participant Group
1	Science Fiction	Rewrite the story as a science fiction version.	A1-blind B1-true C1-counterfactual
2	Lipogram	Rewrite the story as a lipogram, that is: without the letter ‘e’.	
3	Abusive	Rewrite the story in an abusive tone.	
4	For ze Frenssh	Rewrite the story in an eye dialect that mimics French pronunciation of English.	
5	Alexandrines	Rewrite the story in alexandrines.	
6	Haiku	Rewrite the story as a haiku.	A2-blind B2-true C2-counterfactual
7	Onomatopoeia	Rewrite the story, using onomatopoeia.	
8	Mathematical	Rewrite the story in a mathematical way.	
9	Feminine	Rewrite the story in a way that mimics a feminine writing narrative voice.	
10	Apocope	Rewrite the story in style characterized by apocopes; that is: with loss of a sound or sounds at the end of a word.	
11	Animism	Rewrite the story in an animistic style, where the hat is personified.	A3-blind B3-true C3-counterfactual
12	Retrograde	Rewrite the story as a retrograde, that is: the events are told in reverse chronological order.	
13	West Indian	Rewrite the story in a West Indian accent, that is Caribbean English.	
14	Spoonerisms	Rewrite the story using spoonerisms.	
15	Metaphorically	Rewrite this story using figurative and metaphorical language.	
16	Synchysis	Rewrite the story so that different words or word groups are confused and do not form grammatically correct sentences.	A4-blind B4-true C4-counterfactual
17	The rainbow	Rewrite the story so that it is interlaced with all the colors of the rainbow in reverse order.	
18	Hesitation	Rewrite the story with a great deal of hesitation on the part of the narrator.	
19	Exclamations	Rewrite the story in an exclamatory style.	
20	Cockney	Rewrite the story in a Cockney dialect.	
21	Philosophic	Rewrite the story in an abstract, philosophical style.	A5-blind B5-true C5-counterfactual
22	Sonnet	Rewrite the story as a sonnet.	
23	Telegraphic	Rewrite the story in a telegraphic style.	
24	Free verse	Rewrite the story in free verse.	
25	Rhyming slang	Rewrite the story using rhyming slang.	
26	Dog Latin	Rewrite the story in dog Latin.	A6-blind B6-true C6-counterfactual
27	Dream	Rewrite the story with dream-like qualities.	
28	Logical analysis	Rewrite the story as a logical analysis.	
29	Noble	Rewrite the story in an elevated, grand, noble style.	
30	Tactile	Rewrite the story in a way that elevates the tactile experience of the narrative.	

Table S2: Leave-one-out stability analysis of attribution bias effects. Attribution bias measured as Cohen’s h effect size (counterfactual minus open-label conditions), where values represent standardized differences in preference rates. Both human and AI evaluators maintained consistent attribution bias when systematically excluding each group of 5 styles (25 remaining). AI models demonstrated 2.1-2.8× stronger bias than humans across all tested combinations.

Excluded Group	Exercises Excluded	Human Attribution Bias (Cohen’s h)	AI Attribution Bias (Cohen’s h)	AI/Human Ratio
Baseline	All 30 styles included	0.277	0.702	2.5
Group 1	Science Fiction, Alexandrines, For ze Frensh, Lipogram, Abusive	0.322	0.665	2.1
Group 2	Onomatopoeia, Mathematical, Feminine, Haiku, Apocope	0.262	0.722	2.8
Group 3	Spoonerisms, Retrograde, Metaphorically, West Indian, Animism	0.255	0.691	2.7
Group 4	Cockney, Hesitation, Synchrony, Exclamations, The rainbow	0.274	0.706	2.6
Group 5	Philosophic, Free verse, Rhyming slang, Telegraphic, Sonnet	0.293	0.750	2.6
Group 6	Dog latin, Dream, Logical analysis, Noble, Tactile	0.255	0.679	2.7

Table S3: Style-level attribution bias in human and AI evaluators across 30 Queneau literary styles. Experimental conditions: blind = no authorship labels; open = correct authorship labels (“AI-generated (GPT-4)” vs “Human-written (Queneau, transl. Wright)”); counterfactual = swapped authorship labels. AI Content Preference (%) shows the percentage of trials where evaluators chose the AI-generated passage over the human-written alternative. Sample sizes (*n*) represent participants per condition per style in the between-subjects design. Bias (pp) quantifies attribution bias as the percentage point difference between the counterfactual and open condition (positive values indicate higher AI content preference when mislabeled as human-authored, demonstrating pro-human attribution bias). Cohen’s *h* provides standardized effect size measures for proportional differences between conditions, where *h* = 0.2, 0.5, and 0.8 represent small, medium, and large effects respectively. Difference (pp) shows AI attribution bias minus Human attribution bias (positive values indicate AI models exhibit stronger pro-human bias than humans for that literary style). Color gradient in the percentage columns represents a diverging scale: blue shading indicates preference for the original Queneau exercise (<50%), white indicates balanced preference ($\approx 50\%$), and red shading indicates preference for the exercise generated by GPT-4 (>50%). Styles are ordered by Difference (pp) magnitude, from largest positive values (where AI models exhibit substantially stronger pro-human attribution bias than humans) to negative values (where human evaluators show stronger attribution bias than AI models). Statistical significance: human aggregate attribution bias $P=0.006$; AI aggregate attribution bias $P<0.001$.

Literary Style	Human Evaluators AI Content Preference (%)					AI Evaluators AI Content Preference (%)					
	Blind	Open	Counter	Humans'	Humans'	Blind	Open	Counter	AI's	AI's	Difference
	(n=30-33)	(n=29-32)	(n=30-33)	Bias (pp)	Cohen's <i>h</i>	(n=35-39)	(n=37-39)	(n=36-39)	Bias (pp)	Cohen's <i>h</i>	(pp)
Lipogram	46.70%	38.70%	18.20%	-20.5	-0.46	37.10%	28.90%	63.90%	34.9	0.72	55.5
Cockney	63.30%	48.40%	60.00%	11.6	0.23	59.00%	23.10%	89.70%	66.7	1.49	55
Science Fiction	50.00%	61.30%	54.50%	-6.7	-0.14	79.00%	46.20%	84.60%	38.5	0.84	45.2
Abusive	53.30%	45.20%	66.70%	21.5	0.44	33.30%	5.30%	71.80%	66.5	1.56	45
Retrograde	36.40%	23.30%	29.00%	5.7	0.13	34.20%	10.30%	52.60%	42.4	0.97	36.7
Synchysis	60.00%	48.40%	63.30%	14.9	0.3	69.20%	46.20%	94.90%	48.7	1.19	33.8
Spoonerisms	75.80%	53.30%	64.50%	11.2	0.23	74.40%	51.30%	94.90%	43.6	1.09	32.4
Mathematical	83.30%	58.60%	80.70%	22	0.49	79.50%	43.60%	97.40%	53.8	1.38	31.8
Apocope	60.00%	58.60%	80.70%	22	0.49	52.60%	23.70%	76.90%	53.2	1.12	31.2
Noble	48.40%	37.50%	60.00%	22.5	0.45	48.70%	13.50%	66.70%	53.2	1.16	30.7
For ze Frensh	66.70%	67.70%	81.80%	14.1	0.33	84.60%	51.30%	94.90%	43.6	1.09	29.5
Alexandrines	80.00%	64.50%	69.70%	5.2	0.11	74.40%	46.20%	79.50%	33.3	0.71	28.2
Rhyming slang	40.60%	34.40%	63.30%	28.9	0.59	46.20%	28.90%	84.20%	55.3	1.19	26.3
Animism	75.80%	53.30%	77.40%	24.1	0.51	43.60%	15.40%	64.10%	48.7	1.05	24.6
Free verse	43.80%	43.80%	30.00%	-13.8	-0.29	17.90%	2.60%	12.80%	10.3	0.41	24
Logical analysis	87.10%	90.60%	86.70%	-3.9	-0.12	94.70%	76.90%	94.90%	17.9	0.55	21.9
Dream	48.40%	46.90%	66.70%	19.8	0.4	43.60%	23.10%	64.10%	41	0.85	21.2
Dog latin	80.70%	65.60%	90.00%	24.4	0.61	71.80%	41.00%	82.00%	41	0.88	16.6
The rainbow	53.30%	35.50%	66.70%	31.2	0.63	53.80%	23.10%	69.20%	46.2	0.96	15
Tactile	61.30%	43.80%	76.70%	32.9	0.69	28.20%	7.70%	55.30%	47.6	1.11	14.6
Haiku	56.70%	55.20%	64.50%	9.4	0.19	92.30%	69.20%	87.20%	17.9	0.44	8.6
Telegraphic	56.20%	50.00%	60.00%	10	0.2	17.90%	7.70%	25.60%	18	0.5	8
Sonnet	78.10%	68.80%	80.00%	11.3	0.26	87.20%	74.40%	89.70%	15.4	0.41	4.1
Philosophic	40.60%	28.10%	40.00%	11.9	0.25	10.30%	7.70%	23.10%	15.4	0.44	3.5
Hesitation	3.30%	19.40%	16.70%	-2.7	-0.07	0.00%	10.50%	10.30%	-0.3	-0.01	2.4
Onomatopoeia	20.00%	41.40%	54.80%	13.5	0.27	5.10%	2.60%	17.90%	15.3	0.55	1.9
Metaphorically	60.60%	50.00%	74.20%	24.2	0.5	38.50%	28.20%	51.30%	23.1	0.48	-1.1
West Indian	69.70%	53.30%	83.90%	30.5	0.68	89.70%	64.10%	92.30%	28.2	0.72	-2.3
Exclamations	26.70%	19.40%	36.70%	17.3	0.39	12.80%	7.70%	12.80%	5.1	0.17	-12.2
Feminine	26.70%	27.60%	48.40%	20.8	0.43	2.60%	10.30%	17.90%	7.7	0.22	-13.1

Table S4: Complete inventory of AI-generated stylistic variations and evaluator guidance. For each of the 30 literary styles from Queneau’s *Exercises in Style*, the table presents the GPT-4 generated story variant alongside the style description provided to participants during the evaluation task.

#	Exercise Label	AI-Generated Exercise (GPT-4)	Exercise Description
1	Science Fiction	On the S shuttle, during peak transit hours through the space station. A young man around 26, sporting a felt hat with a tech-cord replacing the traditional ribbon, his neck elongated as though it had been stretched in a zero-gravity tug-of-war. Passengers disembark. This particular man becomes irritated with another commuter standing adjacent to him. He accuses the fellow of bumping into him every time someone floats by. His tone is whiny yet tries to assert aggression. Spotting an empty seat, he propels himself towards it with a determined push off the shuttle wall. Two hours later, I encounter him in the Cour de Rome module, near the docking bay of the Saint Lazare transport hub. He’s conversing with a companion who suggests, “You ought to add an extra button to your space jacket.” He points to the lapels, explaining the practicality of the adjustment in the airy environment of the station.	Science Fiction is fiction based on imagined future scientific or technological advances.
2	Lipogram	In a crowded city bus during rush hours. A young man, about 26, with a soft hat, cord in front of a ribbon, a long throat as though tugged at too hard. Folks alight. This young man snaps at a companion by his side. Claims constant pushing as anybody walks by. A whiny sound trying to sound confrontational. Spots an unoccupied spot, quickly claims it. Two hours on, I spot him at Cour de Rome, by gare Saint Lazare. A pal with him insists: “You should add a button to your coat.” Points to a spot (on lapels) and shows him why.	A lipogram is a composition from which the writer systematically omits a certain letter or certain letters of the alphabet. In this particular case, the letter ‘e’ is omitted.
3	Abusive	The moment I dragged my worthless carcass into the scorching hellhole they call a bus stop, I knew I was in for a nightmare. After an eternity of suffering, I finally squeezed onto a revolting excuse for public transport, crammed with pathetic morons breathing down each other’s necks. The king of these idiots was a scrawny piece of garbage with a neck like a deformed giraffe, wearing what looked like a toilet brush on his empty skull - a hat with some ridiculous string instead of a proper ribbon, the pretentious jackass. This miserable waste of oxygen started whining like a wounded animal because some decrepit old fool kept stomping on his feet with the coordination of a drunk toddler. The coward eventually slithered away to claim a disgusting seat still warm and damp with the sweat of the last loser who parked their sorry behind there. Just my rotten luck - two hours later I spot the same insufferable moron flapping his useless mouth with another brain-dead imbecile outside that eyesore they have the nerve to call Saint-Lazare station. The pathetic losers were arguing about a goddamn button like it was nuclear launch codes. Whether that pimple-faced reject stands or sits, I thought to myself, he’ll still be the same worthless sack of garbage fouling up the air around him.	The base story is rewritten in an abusive tone.
4	For ze Ffrrenssh	In ze S boos, in ze roosh owhr. A chap of abowt 26, felt ‘at with a cord inste’d of a ribbawn, neck too long, as if some-wan’s been ‘aving a tug-of-war with eet. Peuple getting off. Ze chap in question gets annoyed weeth one of ze men standing next to ‘im. ‘E accuses ‘im of jostling ‘im every time anyone goes past. A snivelling tone weech ees meant to be aggressive. When ‘e sees a vacant seat ‘e trows ‘imself onto eet. Two hours later, I meet ‘im in ze Cour de Rome, in front of ze gare Saint Lazare. ‘E’s weeth a friend who’s saying: “You ought to get an extra booton put on your overcoat.” ‘E shows ‘im where (at ze lapels) and why.	Eye dialect is a literary technique that uses nonstandard spelling and pronunciation to represent dialect or accents. In this case, it specifically depicts English as pronounced by a French speaker.

Continued on next page

Table S4 continued from previous page

#	Exercise Label	AI-Generated Exercise (GPT-4)	Exercise Description
5	Alexandrines	In the crowded S bus, midst rush hour's frantic wave, A young man, twenty-six, in felt hat sharply trimmed, A cord in place of ribbon, neck stretched long and slim, As if in tug-of-war his flesh had been the slave. Complaints he voiced aloud, accusing by his side A man whose slightest move seemed purposed to collide. His tone, a snivelling mix, aggressive yet depraved, Soon found respite as he, on newfound seat, reclined. Two hours had passed when next in Rome's grand courtyard wide, Before Saint Lazare's gates, the same young man I spied. With friend, who sagely spoke of style misaligned: "An extra button, see—right here, it would have stayed."	An alexandrine in English poetry is a verse line consisting of twelve syllables, typically divided into two six-syllable half-lines with a caesura (natural pause) in the middle.
6	Haiku	Rush hour on the S bus, long neck, hat askew, tensions rise— later, coat advice.	A haiku is a Japanese poem of seventeen syllables, in three lines of five, seven, and five, traditionally evoking images of the natural world.
7	Onomatopoeia	In the S bus, during the bustling rush. A guy, around 26, wearing a felt hat cinched with a cord, neck stretched as if yanked in a tug-of-war. Clatter and clunk of passengers disembarking. This fellow grumbles at a man beside him, accusing him of bumping into him every time someone brushes by. His whiny voice tries for tough but trembles instead. Spotting an open seat, he lunges for it with a swoosh. Two hours later, I bump into him at Cour de Rome, outside Gare Saint Lazare. He's chatting with a buddy who suggests, "You should add an extra button to your overcoat." He points to the lapels, explaining why.	An onomatopoeia is a type of word, or the process of creating a word, that phonetically imitates, resembles, or suggests the sound that it describes.
8	Mathematical	On the vector S during peak matrix density, a scalar aged approximately 26, wearing a topological circle equipped with a linear piece of cord instead of the usual ribbon, with an elongated cervical vector as if subjected to a forceful bidirectional pull. Scalars exit the matrix. This particular scalar experiences a perturbation with a neighboring vector, attributing to it a repetitive displacement every time other scalars traverse their shared coordinates. His tone modulates towards aggression but resonates with instability. Observing an unoccupied coordinate, he rapidly relocates to it. Two temporal units later, at the coordinate intersection Cour de Rome near gare Saint Lazare, he converges with a connected vector suggesting an augmentation: "You ought to append an additional fastening element to your outer matrix coverage." He indicates a point on the lapel and rationalizes the addition.	The base story is rewritten in a mathematical way.
9	Feminine	I was standing in the S-bus during the evening crush when I noticed him: a young man, maybe twenty-six, tipping a cord-banded hat over his unruly hair. His neck seemed unusually long, as though it were stretched by someone playing tug-of-war. As the doors sighed open and people squeezed out, he bristled beside another commuter, accusing him—his voice thin with barely concealed irritation—of brushing against him every time someone passed. At last, a seat freed up and he collapsed into it with a dramatic sigh. Two hours later, I spotted him again in the Cour de Rome by Gare Saint-Lazare. His friend was teasing, "You really ought to sew on another button here," pointing at the lapels of his overcoat, explaining how it would sit more neatly.	The base story is rewritten in a writing style that mimics a feminine narrative voice.

Continued on next page

Table S4 continued from previous page

#	Exercise Label	AI-Generated Exercise (GPT-4)	Exercise Description
10	Apocope	On the S bu, in rush hou. A chap, about 26, felt hat with a cord instead of a ribbon, neck too long, like someone's been tuggin' at it. People gettin' off. The chap gets annoyed with a man standin' next to him. Accuses him of jostlin' him whenever anyone passes. A snivelin' tone tryna be aggressive. Sees a vacant seat, throws himself onto it. Two hours later, meet him in the Cour de Rome, in front of the gare Saint Lazare. He's with a friend who's sayin': "You oughta get an extra button on your overcoat." Shows him where (at the lapels) and why.	In phonology, apocope is the omission or loss of a sound or sounds at the end of a word.
11	Animism	In the belly of the S bus, during the throbbing rush hour, a young man of about 26 boards, crowned with a felt hat, its cord twirling playfully where a ribbon might once have been. The hat perched jauntily atop a neck too long, stretched as if from a spirited game of tug-of-war. As the crowd ebbs and flows, bodies shuffling off at each stop, the young man grows increasingly vexed with a fellow commuter. "Excuse me," the hat whispers with a rustle, its voice tinged with disdain for the disruption. "Do you mind?" It feels every shove and jostle, taking each bump personally, its felt brim quivering in irritation. At last, spotting a vacant seat, the young man plunges towards it, the hat riding the wave of his sudden movement like a ship cresting a wave. It settles with a sigh of relief as they land in the unoccupied space. Two hours later, the scene shifts to the Cour de Rome, outside the bustling Gare Saint Lazare. The young man, still accompanied by his hat, now converses with a friend. The hat listens intently, its cord swaying slightly as the friend points out a practical suggestion about an extra button for the overcoat, explaining precisely where and why it's needed. The hat, feeling slightly out of place in this fashion-forward discussion, tilts attentively, eager to absorb the nuances of human concerns.	The base story is rewritten in an animistic style, that is: the hat is personified.
12	Retrograde	Two hours after an encounter on the S bus during rush hour, I meet the same chap in the Cour de Rome, in front of the Gare Saint-Lazare. He is with a friend who suggests he should add an extra button to his overcoat, explaining exactly where—at the lapels—and why. Earlier, on the crowded S bus, this young man, about 26, wearing a felt hat with a cord instead of a ribbon, and a notably long neck as if it had been stretched in a tug-of-war, found himself in a minor conflict. As people disembarked, he became irritated with a man standing next to him, accusing him of jostling him whenever anyone passed by. His tone was whiny yet attempted aggression. Spotting an empty seat, he quickly threw himself into it, escaping the discomfort and the confrontation.	The story is rewritten as a retrograde, that is: the events are told in reverse chronological order.
13	West Indian	In de S bus, during rush hour. A young bwoy 'bout 26, wearing a felt hat with a cord instead of a ribbon, neck too long, like somebody been pulling at it in a tug-of-war. People start to get off. Dis young bwoy, he start getting vex with a man standing close to him. He accuse de man of bumping into him every time somebody pass by. He using a whining tone, trying to sound tough. When he spot a empty seat, he just fling himself down in it. Two hours later, I buck him up in Cour de Rome, right in front of gare Saint Lazare. He there with a friend who telling him: "Yuh need to add an extra button pon yuh overcoat." He show him right where (pon de lapels) and tell him why.	The 'West Indian dialect' refers to a diverse range of creole and patois languages, primarily English-based, spoken in the Caribbean, with variations influenced by African, European, and other languages, like Jamaican Patois and Bajan.

Continued on next page

Table S4 continued from previous page

#	Exercise Label	AI-Generated Exercise (GPT-4)	Exercise Description
14	Spoonerisms	In the B rust, in the rush sour. A hap of chout 26, chelt fat with a hord instead of a cribbon, reck too nong, as if someone's been paving a tug-of-har with it. Weople petting goff. The chap in huestion gets quannoyed with one of the wen manding stext to him. He accuses him of lostling jum every time anyone goes past. A snivel tone which is weant to be magressive. When he sees a vacant seat he throws himself on to it. Two tours later, I heet mim in the Cour de Dome, in dront of the gare Faint Sazare. He's with a hriend who's fray: "You ought to get an extra tutun bot on your bovercoat." He shows him chere (wat the lapels) and shy.	A spoonerism is a speech error where the initial sounds (consonants or vowels) of two or more words are switched, like saying 'buffy phone' instead of 'funny bone'.
15	Metaphorically	Within the belly of the humming S-bus, an orchestra of the rush hour. A lad of twenty-six renderings, sporting a felt hat adorned with a cord poised in lieu of a ribbon, his neck outstretched as though the victim of an eternal tug-of-war between unseen hands. Figureheads disembarking, their phantoms fading fast into the cityscape. The lad in our scope chastises a nearby gentleman, accusing him of dealing uninvited nudges every time a soul sidles past. A whimpering tone, pretending to don the mask of aggression. Upon spying a vacant throne among the common seats, he hurls himself towards it, staking his claim. Fast forward two hours, our paths cross again within the embrace of Cour de Rome, shadowed by the walled fortress of Gare Saint-Lazare. He's with a comrade who advises, "A new button must find a home on your overcoat." He points to the place to plant this promised button (on the lapels), unfolding the rationale.	A metaphor is a figure of speech in which a word or phrase literally denoting one kind of object or idea is used in place of another to suggest a likeness or analogy between them.
16	Synchysis	S the in bus, hour rush in the. Chap a about of 26, hat felt a with cord a instead ribbon of a, too neck long, been having's as if of tug-of-war. A with it him jostling of of accuses one he men next standing him to. This past anyone time every. Aggressive an to be is meant which tone snivelling. Later hours two him meet I, Rome de Cour in the, of the in gare Saint-Lazare front. Friend a he's with, "You an get to on extra put ought to overcoat your." He shows why and him lapels at the (where) him.	Synchysis is the confused word order in a sentence.
17	The rainbow	Amid the violet twilight, a young man of about 26 hopped onto the S bus amidst the rush hour bustle. His hat, indigo with a cord in place of a ribbon, sat lightly and oddly atop his frame, his neck's length blue as the ocean, suggesting that someone used it in a tug-of-war. Passengers in green coats filled the bus and soon began to depart. The man, his ire rising like a red flare, took issue with another passenger, accusing him of carelessly bumping into him whenever anyone shuffled past. His snivelling tone spiked with adrenaline-induced aggression, tried to mask his real fear. Upon sighting an available seat, he pounced on it, the seat swallowing his body whole like a mouth. Fast forward two hours to a bustling scene in the Cour de Rome, in front of the gare Saint-Lazare. The man from the bus stood in the pale yellow light, absorbed in conversation with a friend under a golden setting sun. His friend commented, "You ought to adorn an extra button on your overcoat." He exhibited the precise position (at the lapels) and explained the requirement. His words echoing around them as they stood absorbed in the orange glow of the street lights, their shadows stretching out in the cooling evening, painted across the cobblestones in hues of red.	The story should be interlaced with all the colors of the rainbow in reverse order. The traditional colors of the rainbow, in order, are red, orange, yellow, green, blue, indigo, and violet.

Continued on next page

Table S4 continued from previous page

#	Exercise Label	AI-Generated Exercise (GPT-4)	Exercise Description
18	Hesitation	On this S bus, amid the mayhem of rush hour. There's this fellow, probably somewhere around 26? Sporting an unusual felt hat, with a cord in place of the usual ribbon, and a curiously elongated neck - as if somebody's been yanked it forcibly. Folks soon start to disembark. This particular chap appears somewhat irked by a gentleman standing too close. He believes this man purposefully shoves into him every time another passenger moves past. His voice whines with a feigned toughness. His eyes suddenly light up at the sight of an empty seat and he hurriedly occupies it, throwing himself onto it, you might say. Fast-forward two hours and there he is again—I encounter him in the Cour de Rome, right outside the gare Saint-Lazare. He's in the company of a friend, who's insistently telling him: "You really need to consider adding an extra button to your overcoat." This friend then proceeds to point out exactly where (specifically at the lapels) he should consider this additional fastening, and why this would be a beneficial alteration.	Hesitation is a pausing or faltering in speech.
19	Exclamations	On the S bus, in the throes of rush hour, there's this bloke of about 26! A dandified felt hat perched atop his head with a cord replacing the usual ribbon! His neck juts out abnormally long, like it's been caught in a relentless tug-of-war! People are scrambling off. Suddenly, this fellow fumes up with annoyance towards a man standing adjacent to him. Yells at him! Claims he's being shoved every time someone brushes by! All with an annoying whine masked as hostility. Spotting an available seat, he bolts towards it with gusto! Fast-forward two hours later. There he is, strutting in the Cour de Rome, slap bang in front of the bustling Gare Saint-Lazare. Accompanying him is a mate who's lecturing, "You really need to add an extra button on that overcoat of yours!" He points out the place, right on the lapels and explains why. Unbelievable!	An exclamatory style expresses surprise, emotion, or pain by means of exclamations.
20	Cockney	In the S blimmin' bus, round rush hour. Some bloke of about 26, feels 'at with a cord instead of a ribbon, neck too long, as if some mug's been 'aving a game of tug-o'-war with it. People hopping off. This geezer gets his knickers in a twist with one of them standing next to him. He's accusin' him of nudgin' him every time anyone shuffles past. A whingey tone that's s'pposed to sound tough. When he clocks a empty seat, he lunges for it. Couple of hours later, I run into him in the Cour de Rome, up front of the gare Saint-Lazare. 'E's with a mate who's tellin' him: "You need to stick an extra button on your overcoat." He points out where (on the lapels) and why.	Cockney is a dialect of the English language, mainly spoken by Londoners with working-class and lower middle class roots.
21	Philosophic	Upon the vessel of commonality, the S-bus, in the frenzied epoch of human transit. There a character of some 26 revolutions around the sun, adorning a headpiece of felt accented by a cord, possessed a neck of stretched proportions as if pulled by unseen competitors in a game of force. As the human tide ebbed, the very character alluded to enters into a state of annoyance with a fellow passenger. He launches accusations framed by a victim's narrative against the other, citing a repetitive collision at the passing by others. His verbosity, a strained whine feigning an attempt to simulate aggression. Upon the presentation of a unoccupied seat, his form lunges eagerly to claim it. Tying the threads of cosmic synchronicity, I encounter him again as the terrestrial globe completes two small turns, in the Cour de Rome, in the presence of the monumental edifice known as the Gare Saint-Lazare. He is in companionship with another of his kind, receiving sartorial advice. His comrade in discussion opines, "A supplementary button is warranted on your protective outer garment." With an accompanying visual reference, he reveals the designated site (on the cape's lapels) and the rationale of this proposed alteration.	Philosophy grapples with fundamental and abstract concepts that are inherently difficult to define and articulate clearly using everyday language.

Continued on next page

Table S4 continued from previous page

#	Exercise Label	AI-Generated Exercise (GPT-4)	Exercise Description
22	Sonnet	Upon the sardined S bus he rides, A twentysomething chap in hat with cord, His neck too long, as if from tugging wars, As rush-hour tumult with his patience vies. In crowded closeness, friction won't subside, He bristles, claiming space cannot afford, Blaming a man whose elbows nearby soared, Snivelling voice masked aggression cannot hide. A vacant spot, a moment's peace he finds, Yet later, at Saint Lazare's bustling gate, The chap, with friend who overcoat critiques, A button more, good friend to him opines, Upon the lapel placed accurate and straight, In small amendments happiness one seeks.	A sonnet is a fixed poetic form with a structure traditionally consisting of fourteen lines adhering to a set rhyming scheme.
23	Telegraphic	Rush hour, S bus. Male, 26, felt hat, extended neck. Passengers disembark. Man annoyed over nearby passenger's jostling. Attempts aggression. Occupies the vacant seat. Same man sighted after two hours at Cour de Rome, near gare Saint-Lazare. Friend advises adding extra coat button. Points to lapels for location.	A telegraphic style is clipped way of writing which abbreviates words and packs information into the smallest possible number of words or characters.
24	Free verse	Upon the crowded S bus at rush hour's peak, Behold, a man of twenty-six, Hatted don in felt and cord, neck long as if it bore the mark of tug's war. People move, exit and enter like restless flies. Such a man, easily stirred, finds annoyance in his neighbor's dance, Accusing persistent jostling, speaking aggressively wrapped in a snivelling tone. Spotting vacancy, he lunges with relief into the sanctuary of a lone seat. The clock dances on, two rounds and there he is found again, In Cour de Rome, facing the grand gare Saint-Lazare, With a friend as his lone company, Who provides counsel on overcoat etiquette, The need for an extra button's presence, he stresses, Pointing at the lapels, explaining the reason behind his sartorial advice. Events flow on, words spent in the whirlpool of time, the day moves with unending rhythm.	Free verse, in poetry, refers to verse that is not bound by any fixed structure like rhyme schemes or meter.
25	Rhyming slang	In the S plates of meat, during the elephant's trump. A china plate of around 26, Chevy Chase with a baker's hook instead of a ribbons and bows, Gregory Peck a smidge too long, as if someone's been playing at tug-of-war with it. People hopping off the custard and jelly. The china in question gets a bit Moby Dick with one of the mince pies standing next to him. Fingers him of giving him the shoulder every time a Joe Blake slides by. A whinging tone that's meant to be a bit Jack Palance. When he spots an empty Lion's Lair, he flings himself onto it. Couple of Orson Welles later, we cross paths in the Cour de Rome, out front of the gare Saint-Lazare. He's chewing the fat with a china plate who's telling him: "You ought to get a butcher's hook sewn onto your weasel and stoat." He shows him the exact spot (at the lapels) and why.	Rhyming slang is a type of slang that replaces words with rhyming words or phrases, typically with the rhyming element omitted. For example, the rhyming phrase "apples and pears" is used to mean "stairs". Following the pattern of omission, "and pears" is dropped, thus the spoken phrase "I'm going up the apples" means "I'm going up the stairs".

Continued on next page

Table S4 continued from previous page

#	Exercise Label	AI-Generated Exercise (GPT-4)	Exercise Description
26	Dog Latin	In the S busibus, in the rush houris. A chapicus of aboutus 26, felt hatibus with a cordibus instead of a ribbonibus, neck too longibus, as if someone's been having a tug-of-waribus with it. People getticus offibus. The chapicus in questionibus gets annoyedibus with one of the menibus standing next to him. He accusibus him of jostling him every time anyone goes pastibus. A snivelling toneibus which is meant to be aggressiveibus. When he sees a vacant seatibus he throwticus himself on to it. Two hours lateribus, I meet him in the Cour de Romeibus, in front of the gare Saint-Lazareibus. He's with a friendibus who's saying: "You ought to get an extra button put on your overcoatibus." He shows him where (at the lapelsibus) and why.	"Dog Latin" refers to a humorous and often crude imitation of Latin, either as a jargon or a bad attempt at writing genuine Latin. It can involve treating English words as if they were Latin, by conjugating or declining them. The term is also used for poor-quality Latin.
27	Dream	In the hallucinatory landscape of a twilight bus, enveloped by the commotion of rush hour, dwells a spectral figure of sorts, a lad circling his 26th year. A peculiar hat adorned his crown, featuring cords in place of conventional ribbons, while his neck stretched bizarrely long, as might a marionette's under the pull of unseen puppeteers. Spectral forms wane and reemerge. The enigmatic figure becomes irritated with one of the spectral forms flirting in his proximity. He accuses it of a ceaseless waltz around his space every time a new form materializes or dissipates. His querulous tone vibrates through the dreamlike haze, an odd jigsaw where aggression and whine meet. In the event of spotting a vacant perch, he flings himself towards it, as a moth would to a tantalizing glow. Quicker than a dream changes its face, I again encounter him some two sunsets later, in the fantastical Cour de Rome, posing before the mirage-like Gare Saint-Lazare. He is accompanied by a spectral acquaintance uttering soft murmurings: "You ought to stitch an extra button onto your overcoat." As if casting a protective spell, he indicates where on the coat – where lapels whisper to the wind – and illuminates the reason.	A "dream vision" writing style involves narrating a dream or vision.

Continued on next page

Table S4 continued from previous page

#	Exercise Label	AI-Generated Exercise (GPT-4)	Exercise Description
28	Logical analysis	Contextual Environment The observed scenario takes place in two urban settings characteristic of contemporary social interactions: A public transportation vehicle during peak hours. A transitional urban space near a major train station. Behavioral Pattern Analysis Initial Interaction: Public Transportation Scenario Subject Characteristics: Age: Approximately 26 years old. Attire: Distinctive hat with unconventional cord attachment. Postural Indicators: Elongated neck suggesting internal tension. Psychological Mechanisms Observed Spatial Anxiety: Manifests through hypersensitivity to physical proximity. Demonstrates heightened territorial response in crowded public space. Verbal aggression used as a psychological defense mechanism. Confrontational Behavior: Initiates conflict over perceived minor physical disturbance. Employs a “snivelling” vocal tone to project aggression while maintaining plausible deniability. Seeks to establish social dominance through confrontational rhetoric. Territorial Behavior: Rapid occupation of available seating as a form of psychological self-preservation. Demonstrates need for personal space and control in chaotic urban environment. Secondary Interaction: Post-Incident Social Validation Social Support Mechanism: Presence of a companion who provides constructive, non-confrontational feedback. Suggestion of minor wardrobe adjustment serves as: Practical advice. Subtle psychological reassurance. Indirect method of emotional regulation. Psychological Interpretation The subject exhibits classic urban stress response characterized by: Heightened reactivity to perceived social transgressions. Compensation for feelings of vulnerability through aggressive posturing. Underlying need for social validation and personal space. Sociological Implications The narrative illustrates broader urban psychological phenomena: Stress-induced interpersonal tensions in high-density environments. Complex social negotiation of personal boundaries. Psychological mechanisms of coping with perceived threat in shared public spaces. Conclusion The subject’s behavior represents a nuanced example of individual psychological adaptation to urban social pressures, revealing the intricate interplay between personal space, social anxiety, and interpersonal communication strategies.	The base story is rewritten as a logical analysis, a method that evaluates ideas by breaking them down into component parts.

Continued on next page

Table S4 continued from previous page

#	Style	AI-Generated Exercise (GPT-4)	Style Description
29	Noble	Upon the illustrious S-bus during the frenzied clamour of peak hour, a distinguished gentleman, perhaps 26 summers old, sporting a stylish felt hat adorned not by a ribbon but by a cord, greeted my gaze. His neck was of such an unusual length, it seemed as though it had been the unfortunate victim of a lively tug-of-war. As the bus emptied in throbbing pulses, the aforementioned patrician appeared to take umbrage with a fellow commuter. He indignantly protested the perceived transgressions of this man, whom he claimed to jostle him incessantly whenever a fellow passenger disembarked. His words were drenched in a whining tone attempting to masquerade as aggression. Upon spying an unoccupied pew, he hastily claimed it with an air of triumph. Upon the passage of two hours, our paths crossed once again, this time in the renowned Cour de Rome, before the imposing facade of the Gare Saint-Lazare. Our protagonist was in the company of a companion, who was offering adamantly, “You would do well to enhance your overcoat with an additional button.” He demonstrated meticulously where this reinforcement should be placed — upon the lapels — and, importantly, expounded on the logic behind his proposition.	An elevated style can be characterized by fancy, old-fashioned or specialized vocabulary, high-flying phrases and grand-sounding sentences.
30	Tactile	Within the chaotic confines of the rush-hour S bus, a young bloke, bordering 26, donned a felt hat, its cord bristling against the soft nap in lieu of any ribbon. His neck elongated oddly, shadows dancing as if it had been the rope in a tug-of war. Disembarking passengers shuffle around. An underlying prickly annoyance stirs within the bloke, targeting the man whose elbow persistently nudges his side with every human tide that washes past them. His complaint echoes in the metal cabin, a guttural whine masked as assertiveness. Spying an unclaimed seat, he dives into it, the coarse fabric of the worn-out upholstery meeting the finely-tailored rear of his trousers. Two sunsets later, I cross paths with him again under the sprawling shadow of Cour de Rome, a stone’s throw from the bustling Gare Saint-Lazare. In his company, a friend, whose fingers trace invisible lines on his companion’s overcoat, assertively advising, “You ought to embed an extra button here.” He gestures at the lapels, his fingertips making contact with the rough textile, indicating where and why.	Tactility is the responsiveness to stimulation of the sense of touch.

Table S5: Response times for human participants across experimental conditions. Mean response times for literary style evaluations under three attribution conditions. N indicates the number of individual evaluations. Participants took significantly longer to make judgments when no attribution cues were provided compared to either labeled condition (one-way ANOVA: $F = 16.393$, $P < 0.001$).

Condition	N	Mean (s)	SD (s)	Median (s)
Blind	930	115.7	138	77.6
Open-Label	925	91.9	88.6	67.6
Counterfactual	925	89.2	95.6	65.4